

# 6\_Backpropagation\_students

July 13, 2019

## 1 HaMLeT

### 1.1 Session 6: Backpropagation

by Leon Weninger and Raphael Kolk

#### 1.1.1 Goal of this Session

In this session you will, step by step, implement a the backpropagation algorithm yourself without using any deep learning libraries. You should already be familiar with Python as well as NumPy (a package for scientific computing with Python).

#### 1.1.2 Given code

**Task 0:** Familiarize yourself briefly with the given code. Pay particural attention to the Layer and Cost classes, from which you will derive the classes you implement, and the Sigmoid layer which is predefined as an example. You'll also need to execute the cells in this section once.

The following code loads the data and trains the network in a similar fashion as in the last session.

```
In [9]: import numpy as np
        from tqdm import tqdm
        from load_mnist import load_mnist

        def vectorize(j):
            label_vector = np.zeros((1, 10))
            label_vector[0, int(j)] = 1.0
            return label_vector

        def load_data():
            images, labels = load_mnist()

            image_size = images.shape[1]
            label_size = labels.shape[1]

            random_permutation = np.random.permutation(images.shape[0])
```

```

images = images[random_permutation, :]
labels = labels[random_permutation, :]

return images, labels, image_size, label_size

def train(net, cost_function, number_epochs, batch_size, learning_rate):
    images, labels, image_size, label_size = load_data()
    training_images, validation_images = images[:50000], images[50000:]
    training_labels, validation_labels = labels[:50000], labels[50000:]

    for e in range(number_epochs):
        cost = train_epoch(e, net, training_images, training_labels, cost_function, batch_size, learning_rate)
        accuracy = validate_epoch(e, net, validation_images, validation_labels, batch_size, learning_rate)
        print('cost=%5.6f, accuracy=%2.6f' % (cost, accuracy), flush=True)

def train_epoch(e, net, images, labels, cost_function, batch_size, learning_rate):
    epoch_cost = 0

    for i in tqdm(range(0, len(images), batch_size), ascii=False, desc='training, e=%d' % e):
        batch_images = images[i:min(i + batch_size, len(images)), :]
        batch_labels = labels[i:min(i + batch_size, len(labels)), :]

        # zero the gradients
        net.zero_gradients()

        # forward pass
        prediction = net.forward(batch_images)
        cost = cost_function.estimate(batch_labels, prediction)

        # backward pass
        dprediction = cost_function.gradient(cost)
        net.backward(dprediction)

        # update the parameters using the computed gradients via stochastic gradient descent
        net.update_parameters(learning_rate)

    epoch_cost += np.mean(cost)

    return epoch_cost

def validate_epoch(e, net, images, labels, batch_size):
    n_correct = 0
    n_total = 0

    for i in tqdm(range(0, len(images), batch_size), ascii=False, desc='validation, e=%d' % e):

```

```

batch_images = images[i:min(i + batch_size, len(images)), :]
batch_labels = labels[i:min(i + batch_size, len(labels)), :]

# compute predicted probabilities.
predictions = net.forward(batch_images)

# find the most probable class label.
n_correct += sum(np.argmax(batch_labels, axis=1) == np.argmax(predictions, axis=1))
n_total += batch_labels.shape[0]

return n_correct / n_total

```

Remember the sigmoid function and its derivative you implemented in the previous session.

```

In [10]: def sigmoid_function(var):
        return 1.0 / (1.0 + np.exp(-var))

def sigmoid_derivative(z):
    return sigmoid_function(z) * (1 - sigmoid_function(z))

```

The following abstract classes should serve as parent classes for all the different layers and cost functions which you will implement.

```

In [11]: class Layer:
        def __init__(self):
            # Initialize all member variables of the layer.
            pass

        def forward(self, x_in):
            # Implements for forward pass of the layer and returns x_out.
            pass

        def backward(self, d_out):
            # Implements the backward pass of the layer and returns d_in.
            pass

        def zero_gradients(self):
            # Sets all gradients of the layer to zero.
            pass

        def update_parameters(self, learning_rate):
            # Update the parameters of the layer with the help of the gradients stored du
            pass

class Cost:
    def __init__(self):
        # Initialize all member variables of the cost function.

```

```

        pass

    def estimate(self, target, prediction):
        # Estimates and return the cost with respect to the predicted label and a tar
        pass

    def gradient(self, cost):
        # Calculates and returns the gradient with respect to the cost.
        pass

```

The following class derived from the Layer class implements the forward and backward pass of the sigmoid activation function already known from the previous session and serves as an example for you. Since it does not have learnable parameters, no `update_parameters` or `zero_gradients` function needs to be implemented.

```

In [12]: class Sigmoid(Layer):
    def __init__(self):
        self.x_in = None

    def forward(self, x_in):
        self.x_in = x_in
        x_out = sigmoid_function(x_in)
        return x_out

    def backward(self, d_out):
        d_in = d_out * sigmoid_derivative(self.x_in)
        return d_in

    def zero_gradients(self):
        pass

    def update_parameters(self, learning_rate):
        pass

```

### 1.1.3 Theoretical Foundation

**Task 1:** Take a piece of paper and a pencil, use your knowledge from the preparation material and the introduction slides and fill in the gaps in the preparation material. Keep in mind that the inputs may be batched. Having written the formulas down, please check with the tutor if they are correct.

### 1.1.4 Practical Implementation

**Task 2a:** Implement the forward function for the Linear layer. Remember to store the `x_in` for use in the backward pass.

**Task 2b:** Implement the `estimate` function for the `MeanSquareError` cost, which estimates the cost after a ground truth target is set. Remember to store the prediction for calculating the gradient.

**Task 2c:** Implement the gradient function for the MeanSquareError cost, which calculates the gradient with respect to the cost. Use the prediction stored during the forward pass.

**Task 2d:** Implement the backward function for the Linear layer. The function should also calculate and accumulate the gradient of  $w$  and  $b$  with regard to the error.

**Task 2e:** Implement the `update_parameters` function for the Linear layer, i.e., use the gradients  $dw$  and  $db$  together with a given `learning_rate` to update the parameters  $w$  and  $b$  accordingly.

**Task 2f:** Test your implementation by propagating random input through a linear layer followed by a sigmoid layer, estimating the mean square error to a random target, calculating the gradient and propagating it back through the sigmoid and linear layer. Afterwards update the parameters of the linear layer using the function you implemented.

**Task 3a:** Implement the `Network` class which can encapsulate multiple layers. It offers the same interface as a layer and is therefore derived from the `Layer` parent as well. Make sure to implement all member functions needed. The forward function propagates a given input through all encapsulated layers and returns the final prediction of the network, whereas the backward function propagates a given gradient through all layers in reversed order. `zero_gradients` and `update_parameters` invoke the respective functions of the encapsulated layers.

**Task 3b:** Test your implementation analogous to task 2f but using the `Network` class to encapsulate the linear and sigmoid layer.

**Task 4:** Train the network you just implemented using the `dataloader` and `train` function given above and the hyperparameter given below.

**Task 5:** Come up with a more sophisticated network structure and adjust the hyperparameter in order to increase the accuracy.

```
In [13]: class Linear(Layer):
        def __init__(self, n_in, n_out, initial_sigma=0.1):
            self.n_in = n_in
            self.n_out = n_out

            self.w = initial_sigma * np.random.randn(n_out, n_in)
            self.b = np.zeros((1, n_out))

            self.zero_gradients()

            self.x_in = None

        def forward(self, x_in):
            # ----- Add code for task 2a between comments -----
            self.x_in = x_in

            x_out = np.dot(self.x_in, self.w.T) + self.b

            # -----
            return x_out

        def backward(self, d_out):
            # ----- Add code for task 2d between comments -----

            self.db += np.sum(d_out, axis=0, keepdims=True)
```

```

        self.dw += np.dot(d_out.T, self.x_in)

        self.d_in = np.dot(d_out, self.w)
        # -----
        return self.d_in

    def zero_gradients(self):
        self.dw = np.zeros((self.n_out, self.n_in))
        self.db = np.zeros((1, self.n_out))
        self.dx = np.empty((0, self.n_in))

    def update_parameters(self, learning_rate):
        # ----- Add code for task 2e between comments -----

        self.w = self.w - (learning_rate * self.dw)
        self.b = self.b - (learning_rate * self.db)
        # -----

class MeanSquareError(Cost):
    def __init__(self):
        self.prediction = None
        self.target = None

    def estimate(self, target, prediction):
        self.target = target
        self.prediction = prediction
        # ----- add code for task 2b between comments -----
        cost = 1/2 * (target - prediction)**2
        # -----
        return cost

    def gradient(self, cost):
        # ----- add code for task 2c between comments -----

        gradient = self.prediction - self.target

        # gradient = np.sqrt(2 * cost)

        # -----
        return gradient

class Network(Layer):
    def __init__(self, layers):
        self.layers = layers

```

```

# ----- add code for task 3a between comments -----
def forward(self,x):
    for layer in self.layers:
        x=layer.forward(x)
    return x

def backward(self,d):
    for layer in reversed(self.layers):
        d=layer.backward(d)
    return d

def update_parameters(self,alpha):
    for layer in self.layers:
        layer.update_parameters(alpha)

def zero_gradients(self):
    for layer in self.layers:
        layer.zero_gradients()

# -----

```

```

In [6]: # define hyperparameters
input_size = 28**2
label_size = 10
batch_size = 600
learning_rate = 0.0001
number_epochs = 100

random_input = np.random.rand(batch_size, input_size)
random_label = np.random.rand(batch_size, label_size)

linear_layer = Linear(input_size, label_size)
sigmoid_layer = Sigmoid()
cost_function = MeanSquareError()

# ----- add code for task 2f between comments -----
act=linear_layer.forward(random_input)
pred=sigmoid_layer.forward(act)

cost=cost_function.estimate(random_label,pred)
grad=cost_function.gradient(cost)

d=sigmoid_layer.backward(grad)
d_in=linear_layer.backward(d)

linear_layer.update_parameters(learning_rate)

# -----

```

```

# ----- add code for task 3b between comments -----

linear_layer_1 = Linear(input_size, label_size)
sigmoid_layer = Sigmoid()
net = Network([linear_layer_1, sigmoid_layer])
train(net, cost_function, number_epochs, batch_size, learning_rate)

# -----

# ----- add code for task 4 between comments -----

# -----

# ----- add code for task 5 between comments -----

# -----

training, e=0: 0%|          | 0/84 [00:00<?, ?it/s]/usr/local/anaconda3/lib/python3.6/site
training, e=0: 100%|| 84/84 [00:00<00:00, 98.33it/s]
validation, e=0: 100%|| 17/17 [00:00<00:00, 185.19it/s]

cost=4.938776, accuracy=0.309500

training, e=1: 100%|| 84/84 [00:00<00:00, 104.64it/s]
validation, e=1: 100%|| 17/17 [00:00<00:00, 185.28it/s]

cost=3.624176, accuracy=0.306500

training, e=2: 100%|| 84/84 [00:00<00:00, 105.60it/s]
validation, e=2: 100%|| 17/17 [00:00<00:00, 185.08it/s]

cost=3.551397, accuracy=0.320900

training, e=3: 100%|| 84/84 [00:00<00:00, 106.68it/s]
validation, e=3: 100%|| 17/17 [00:00<00:00, 185.09it/s]

```



cost=3.511281, accuracy=0.333500

training, e=4: 100%|| 84/84 [00:00<00:00, 105.62it/s]  
validation, e=4: 100%|| 17/17 [00:00<00:00, 185.08it/s]

cost=3.461046, accuracy=0.335800

training, e=5: 100%|| 84/84 [00:00<00:00, 86.33it/s]  
validation, e=5: 100%|| 17/17 [00:00<00:00, 141.62it/s]

cost=3.331616, accuracy=0.396200

training, e=6: 100%|| 84/84 [00:00<00:00, 102.03it/s]  
validation, e=6: 100%|| 17/17 [00:00<00:00, 185.37it/s]

cost=3.168227, accuracy=0.413300

training, e=7: 100%|| 84/84 [00:00<00:00, 107.01it/s]  
validation, e=7: 100%|| 17/17 [00:00<00:00, 185.24it/s]

cost=3.058832, accuracy=0.417100

training, e=8: 100%|| 84/84 [00:00<00:00, 104.91it/s]  
validation, e=8: 100%|| 17/17 [00:00<00:00, 185.32it/s]

cost=2.994231, accuracy=0.431600

training, e=9: 100%|| 84/84 [00:00<00:00, 98.41it/s]  
validation, e=9: 100%|| 17/17 [00:00<00:00, 185.79it/s]

cost=2.963683, accuracy=0.428900

training, e=10: 100%|| 84/84 [00:00<00:00, 108.27it/s]  
validation, e=10: 100%|| 17/17 [00:00<00:00, 185.33it/s]

cost=2.921293, accuracy=0.438800

training, e=11: 100%|| 84/84 [00:00<00:00, 108.18it/s]  
validation, e=11: 100%|| 17/17 [00:00<00:00, 187.51it/s]

cost=2.894430, accuracy=0.441400

training, e=12: 100%|| 84/84 [00:00<00:00, 92.79it/s]  
validation, e=12: 100%|| 17/17 [00:00<00:00, 141.94it/s]

cost=2.871689, accuracy=0.451400

training, e=13: 100%|| 84/84 [00:00<00:00, 108.81it/s]  
validation, e=13: 100%|| 17/17 [00:00<00:00, 184.29it/s]

cost=2.854031, accuracy=0.442300

training, e=14: 100%|| 84/84 [00:00<00:00, 108.94it/s]  
validation, e=14: 100%|| 17/17 [00:00<00:00, 185.11it/s]

cost=2.836971, accuracy=0.443800

training, e=15: 100%|| 84/84 [00:00<00:00, 106.09it/s]  
validation, e=15: 100%|| 17/17 [00:00<00:00, 184.23it/s]

cost=2.827713, accuracy=0.444200

training, e=16: 100%|| 84/84 [00:00<00:00, 92.51it/s]  
validation, e=16: 100%|| 17/17 [00:00<00:00, 139.15it/s]

cost=2.833228, accuracy=0.458200

training, e=17: 100%|| 84/84 [00:00<00:00, 109.96it/s]  
validation, e=17: 100%|| 17/17 [00:00<00:00, 184.68it/s]

cost=2.813443, accuracy=0.459700

training, e=18: 100%|| 84/84 [00:00<00:00, 110.14it/s]  
validation, e=18: 100%|| 17/17 [00:00<00:00, 184.95it/s]

cost=2.804464, accuracy=0.465800

training, e=19: 100%|| 84/84 [00:00<00:00, 110.13it/s]  
validation, e=19: 100%|| 17/17 [00:00<00:00, 185.00it/s]

cost=2.798534, accuracy=0.467900

training, e=20: 100%|| 84/84 [00:00<00:00, 102.17it/s]  
validation, e=20: 100%|| 17/17 [00:00<00:00, 183.10it/s]

cost=2.784008, accuracy=0.468700

training, e=21: 100%|| 84/84 [00:00<00:00, 102.78it/s]  
validation, e=21: 100%|| 17/17 [00:00<00:00, 141.85it/s]

cost=2.779173, accuracy=0.467800

training, e=22: 100%|| 84/84 [00:00<00:00, 100.15it/s]  
validation, e=22: 100%|| 17/17 [00:00<00:00, 185.21it/s]

cost=2.776501, accuracy=0.467600

training, e=23: 100%|| 84/84 [00:00<00:00, 109.24it/s]  
validation, e=23: 100%|| 17/17 [00:00<00:00, 184.87it/s]

cost=2.763745, accuracy=0.469400

training, e=24: 100%|| 84/84 [00:00<00:00, 106.26it/s]  
validation, e=24: 100%|| 17/17 [00:00<00:00, 184.94it/s]

cost=2.756209, accuracy=0.468500

training, e=25: 100%|| 84/84 [00:00<00:00, 85.54it/s]  
validation, e=25: 100%|| 17/17 [00:00<00:00, 181.22it/s]

cost=2.757379, accuracy=0.472200

training, e=26: 100%|| 84/84 [00:00<00:00, 102.88it/s]  
validation, e=26: 100%|| 17/17 [00:00<00:00, 184.70it/s]  
  
cost=2.744545, accuracy=0.472800

training, e=27: 100%|| 84/84 [00:01<00:00, 78.79it/s]  
validation, e=27: 100%|| 17/17 [00:00<00:00, 143.03it/s]  
  
cost=2.740545, accuracy=0.474200

training, e=28: 100%|| 84/84 [00:00<00:00, 109.50it/s]  
validation, e=28: 100%|| 17/17 [00:00<00:00, 179.70it/s]  
  
cost=2.733786, accuracy=0.478800

training, e=29: 100%|| 84/84 [00:00<00:00, 108.12it/s]  
validation, e=29: 100%|| 17/17 [00:00<00:00, 184.78it/s]  
  
cost=2.719958, accuracy=0.479400

training, e=30: 100%|| 84/84 [00:00<00:00, 107.57it/s]  
validation, e=30: 100%|| 17/17 [00:00<00:00, 184.67it/s]  
  
cost=2.707361, accuracy=0.487400

training, e=31: 100%|| 84/84 [00:00<00:00, 106.98it/s]  
validation, e=31: 100%|| 17/17 [00:00<00:00, 185.06it/s]  
  
cost=2.707886, accuracy=0.478700

training, e=32: 100%|| 84/84 [00:00<00:00, 107.55it/s]  
validation, e=32: 100%|| 17/17 [00:00<00:00, 184.66it/s]  
  
cost=2.678279, accuracy=0.485100

training, e=33: 100%|| 84/84 [00:00<00:00, 91.85it/s]  
validation, e=33: 100%|| 17/17 [00:00<00:00, 184.20it/s]

cost=2.666325, accuracy=0.494200

training, e=34: 100%|| 84/84 [00:00<00:00, 87.82it/s]  
validation, e=34: 100%|| 17/17 [00:00<00:00, 147.79it/s]

cost=2.655260, accuracy=0.501600

training, e=35: 100%|| 84/84 [00:00<00:00, 102.87it/s]  
validation, e=35: 100%|| 17/17 [00:00<00:00, 185.91it/s]

cost=2.646725, accuracy=0.497400

training, e=36: 100%|| 84/84 [00:00<00:00, 104.94it/s]  
validation, e=36: 100%|| 17/17 [00:00<00:00, 186.12it/s]

cost=2.630066, accuracy=0.499200

training, e=37: 100%|| 84/84 [00:00<00:00, 103.39it/s]  
validation, e=37: 100%|| 17/17 [00:00<00:00, 171.69it/s]

cost=2.635276, accuracy=0.510300

training, e=38: 100%|| 84/84 [00:00<00:00, 105.63it/s]  
validation, e=38: 100%|| 17/17 [00:00<00:00, 185.57it/s]

cost=2.623842, accuracy=0.537000

training, e=39: 100%|| 84/84 [00:00<00:00, 106.26it/s]  
validation, e=39: 100%|| 17/17 [00:00<00:00, 185.60it/s]

cost=2.581603, accuracy=0.537500

training, e=40: 100%|| 84/84 [00:00<00:00, 104.66it/s]  
validation, e=40: 100%|| 17/17 [00:00<00:00, 185.04it/s]

cost=2.520833, accuracy=0.545400

training, e=41: 100%|| 84/84 [00:00<00:00, 103.92it/s]  
validation, e=41: 100%|| 17/17 [00:00<00:00, 185.01it/s]  
  
cost=2.467117, accuracy=0.550700

training, e=42: 100%|| 84/84 [00:00<00:00, 108.89it/s]  
validation, e=42: 100%|| 17/17 [00:00<00:00, 185.26it/s]  
  
cost=2.397191, accuracy=0.553800

training, e=43: 100%|| 84/84 [00:00<00:00, 109.47it/s]  
validation, e=43: 100%|| 17/17 [00:00<00:00, 185.62it/s]  
  
cost=2.353979, accuracy=0.556800

training, e=44: 100%|| 84/84 [00:00<00:00, 97.42it/s]  
validation, e=44: 100%|| 17/17 [00:00<00:00, 140.45it/s]  
  
cost=2.332697, accuracy=0.558900

training, e=45: 100%|| 84/84 [00:00<00:00, 98.00it/s]  
validation, e=45: 100%|| 17/17 [00:00<00:00, 185.64it/s]  
  
cost=2.311087, accuracy=0.559800

training, e=46: 100%|| 84/84 [00:00<00:00, 110.28it/s]  
validation, e=46: 100%|| 17/17 [00:00<00:00, 185.35it/s]  
  
cost=2.300523, accuracy=0.559800

training, e=47: 100%|| 84/84 [00:00<00:00, 109.90it/s]  
validation, e=47: 100%|| 17/17 [00:00<00:00, 185.31it/s]  
  
cost=2.287708, accuracy=0.561300

training, e=48: 100%|| 84/84 [00:00<00:00, 103.65it/s]  
validation, e=48: 100%|| 17/17 [00:00<00:00, 183.99it/s]

cost=2.282187, accuracy=0.559300

training, e=49: 100%|| 84/84 [00:00<00:00, 102.65it/s]  
validation, e=49: 100%|| 17/17 [00:00<00:00, 140.79it/s]

cost=2.276578, accuracy=0.565600

training, e=50: 100%|| 84/84 [00:00<00:00, 88.83it/s]  
validation, e=50: 100%|| 17/17 [00:00<00:00, 186.06it/s]

cost=2.284204, accuracy=0.563600

training, e=51: 100%|| 84/84 [00:00<00:00, 108.71it/s]  
validation, e=51: 100%|| 17/17 [00:00<00:00, 185.19it/s]

cost=2.280297, accuracy=0.559800

training, e=52: 100%|| 84/84 [00:00<00:00, 109.63it/s]  
validation, e=52: 100%|| 17/17 [00:00<00:00, 185.28it/s]

cost=2.264273, accuracy=0.560100

training, e=53: 100%|| 84/84 [00:00<00:00, 109.92it/s]  
validation, e=53: 100%|| 17/17 [00:00<00:00, 185.38it/s]

cost=2.253747, accuracy=0.562300

training, e=54: 100%|| 84/84 [00:00<00:00, 96.16it/s]  
validation, e=54: 100%|| 17/17 [00:00<00:00, 140.83it/s]

cost=2.235027, accuracy=0.565500

training, e=55: 100%|| 84/84 [00:00<00:00, 100.62it/s]  
validation, e=55: 100%|| 17/17 [00:00<00:00, 185.84it/s]

cost=2.240275, accuracy=0.562500

training, e=56: 100%|| 84/84 [00:00<00:00, 109.86it/s]  
validation, e=56: 100%|| 17/17 [00:00<00:00, 185.31it/s]  
  
cost=2.242104, accuracy=0.563700

training, e=57: 100%|| 84/84 [00:00<00:00, 109.61it/s]  
validation, e=57: 100%|| 17/17 [00:00<00:00, 185.91it/s]  
  
cost=2.230347, accuracy=0.565200

training, e=58: 100%|| 84/84 [00:00<00:00, 109.60it/s]  
validation, e=58: 100%|| 17/17 [00:00<00:00, 185.05it/s]  
  
cost=2.238238, accuracy=0.564800

training, e=59: 100%|| 84/84 [00:00<00:00, 110.51it/s]  
validation, e=59: 100%|| 17/17 [00:00<00:00, 185.30it/s]  
  
cost=2.227793, accuracy=0.566000

training, e=60: 100%|| 84/84 [00:00<00:00, 110.40it/s]  
validation, e=60: 100%|| 17/17 [00:00<00:00, 185.50it/s]  
  
cost=2.220233, accuracy=0.567100

training, e=61: 100%|| 84/84 [00:01<00:00, 68.95it/s]  
validation, e=61: 100%|| 17/17 [00:00<00:00, 132.96it/s]  
  
cost=2.215265, accuracy=0.567500

training, e=62: 100%|| 84/84 [00:00<00:00, 96.49it/s]  
validation, e=62: 100%|| 17/17 [00:00<00:00, 152.10it/s]  
  
cost=2.218362, accuracy=0.567000

training, e=63: 100%|| 84/84 [00:00<00:00, 109.85it/s]  
validation, e=63: 100%|| 17/17 [00:00<00:00, 185.02it/s]



cost=2.209902, accuracy=0.568500

training, e=64: 100%|| 84/84 [00:00<00:00, 109.06it/s]  
validation, e=64: 100%|| 17/17 [00:00<00:00, 184.59it/s]

cost=2.204271, accuracy=0.565200

training, e=65: 100%|| 84/84 [00:00<00:00, 110.31it/s]  
validation, e=65: 100%|| 17/17 [00:00<00:00, 184.42it/s]

cost=2.209475, accuracy=0.570000

training, e=66: 100%|| 84/84 [00:00<00:00, 109.82it/s]  
validation, e=66: 100%|| 17/17 [00:00<00:00, 184.34it/s]

cost=2.203692, accuracy=0.569300

training, e=67: 100%|| 84/84 [00:00<00:00, 109.60it/s]  
validation, e=67: 100%|| 17/17 [00:00<00:00, 184.10it/s]

cost=2.203329, accuracy=0.561700

training, e=68: 100%|| 84/84 [00:00<00:00, 109.12it/s]  
validation, e=68: 100%|| 17/17 [00:00<00:00, 183.43it/s]

cost=2.199000, accuracy=0.572300

training, e=69: 100%|| 84/84 [00:00<00:00, 110.46it/s]  
validation, e=69: 100%|| 17/17 [00:00<00:00, 184.49it/s]

cost=2.188360, accuracy=0.573500

training, e=70: 100%|| 84/84 [00:00<00:00, 110.33it/s]  
validation, e=70: 100%|| 17/17 [00:00<00:00, 184.57it/s]

cost=2.185778, accuracy=0.574800

training, e=71: 100%|| 84/84 [00:00<00:00, 103.58it/s]  
validation, e=71: 100%|| 17/17 [00:00<00:00, 183.74it/s]  
  
cost=2.187600, accuracy=0.567200

training, e=72: 100%|| 84/84 [00:00<00:00, 92.39it/s]  
validation, e=72: 100%|| 17/17 [00:00<00:00, 141.34it/s]  
  
cost=2.195178, accuracy=0.568600

training, e=73: 100%|| 84/84 [00:00<00:00, 110.59it/s]  
validation, e=73: 100%|| 17/17 [00:00<00:00, 184.50it/s]  
  
cost=2.176681, accuracy=0.570600

training, e=74: 100%|| 84/84 [00:00<00:00, 110.14it/s]  
validation, e=74: 100%|| 17/17 [00:00<00:00, 184.73it/s]  
  
cost=2.174270, accuracy=0.571300

training, e=75: 100%|| 84/84 [00:00<00:00, 105.69it/s]  
validation, e=75: 100%|| 17/17 [00:00<00:00, 183.77it/s]  
  
cost=2.175344, accuracy=0.571000

training, e=76: 100%|| 84/84 [00:00<00:00, 106.54it/s]  
validation, e=76: 100%|| 17/17 [00:00<00:00, 184.54it/s]  
  
cost=2.177708, accuracy=0.567100

training, e=77: 100%|| 84/84 [00:00<00:00, 110.35it/s]  
validation, e=77: 100%|| 17/17 [00:00<00:00, 184.66it/s]  
  
cost=2.171814, accuracy=0.573500

training, e=78: 100%|| 84/84 [00:00<00:00, 102.89it/s]  
validation, e=78: 100%|| 17/17 [00:00<00:00, 183.57it/s]

cost=2.172597, accuracy=0.573700

training, e=79: 100%|| 84/84 [00:00<00:00, 104.52it/s]  
validation, e=79: 100%|| 17/17 [00:00<00:00, 184.40it/s]

cost=2.162491, accuracy=0.573200

training, e=80: 100%|| 84/84 [00:00<00:00, 110.31it/s]  
validation, e=80: 100%|| 17/17 [00:00<00:00, 184.57it/s]

cost=2.164888, accuracy=0.572300

training, e=81: 100%|| 84/84 [00:00<00:00, 110.80it/s]  
validation, e=81: 100%|| 17/17 [00:00<00:00, 184.71it/s]

cost=2.167944, accuracy=0.569700

training, e=82: 100%|| 84/84 [00:00<00:00, 110.42it/s]  
validation, e=82: 100%|| 17/17 [00:00<00:00, 184.88it/s]

cost=2.168157, accuracy=0.576600

training, e=83: 100%|| 84/84 [00:00<00:00, 110.62it/s]  
validation, e=83: 100%|| 17/17 [00:00<00:00, 184.48it/s]

cost=2.161013, accuracy=0.573400

training, e=84: 100%|| 84/84 [00:00<00:00, 110.86it/s]  
validation, e=84: 100%|| 17/17 [00:00<00:00, 184.94it/s]

cost=2.155828, accuracy=0.574200

training, e=85: 100%|| 84/84 [00:00<00:00, 111.65it/s]  
validation, e=85: 100%|| 17/17 [00:00<00:00, 184.75it/s]

cost=2.159542, accuracy=0.574700

training, e=86: 100%|| 84/84 [00:00<00:00, 111.03it/s]  
validation, e=86: 100%|| 17/17 [00:00<00:00, 185.27it/s]

cost=2.160014, accuracy=0.572000

training, e=87: 100%|| 84/84 [00:00<00:00, 110.15it/s]  
validation, e=87: 100%|| 17/17 [00:00<00:00, 184.36it/s]

cost=2.157137, accuracy=0.574100

training, e=88: 100%|| 84/84 [00:00<00:00, 109.97it/s]  
validation, e=88: 100%|| 17/17 [00:00<00:00, 184.61it/s]

cost=2.159303, accuracy=0.576700

training, e=89: 100%|| 84/84 [00:00<00:00, 110.28it/s]  
validation, e=89: 100%|| 17/17 [00:00<00:00, 184.68it/s]

cost=2.151820, accuracy=0.573100

training, e=90: 100%|| 84/84 [00:00<00:00, 110.99it/s]  
validation, e=90: 100%|| 17/17 [00:00<00:00, 183.54it/s]

cost=2.152159, accuracy=0.572700

training, e=91: 100%|| 84/84 [00:00<00:00, 110.73it/s]  
validation, e=91: 100%|| 17/17 [00:00<00:00, 184.42it/s]

cost=2.153817, accuracy=0.573500

training, e=92: 100%|| 84/84 [00:00<00:00, 110.49it/s]  
validation, e=92: 100%|| 17/17 [00:00<00:00, 184.44it/s]

cost=2.152125, accuracy=0.573900

training, e=93: 100%|| 84/84 [00:00<00:00, 110.21it/s]  
validation, e=93: 100%|| 17/17 [00:00<00:00, 184.56it/s]

cost=2.152750, accuracy=0.573000

training, e=94: 100%|| 84/84 [00:00<00:00, 110.07it/s]  
validation, e=94: 100%|| 17/17 [00:00<00:00, 184.83it/s]

cost=2.150829, accuracy=0.573100

training, e=95: 100%|| 84/84 [00:00<00:00, 110.35it/s]  
validation, e=95: 100%|| 17/17 [00:00<00:00, 184.69it/s]

cost=2.148718, accuracy=0.575400

training, e=96: 100%|| 84/84 [00:00<00:00, 108.63it/s]  
validation, e=96: 100%|| 17/17 [00:00<00:00, 180.89it/s]

cost=2.143742, accuracy=0.574200

training, e=97: 100%|| 84/84 [00:00<00:00, 108.98it/s]  
validation, e=97: 100%|| 17/17 [00:00<00:00, 180.97it/s]

cost=2.142858, accuracy=0.572100

training, e=98: 100%|| 84/84 [00:00<00:00, 108.55it/s]  
validation, e=98: 100%|| 17/17 [00:00<00:00, 184.65it/s]

cost=2.138705, accuracy=0.571700

training, e=99: 100%|| 84/84 [00:00<00:00, 110.00it/s]  
validation, e=99: 100%|| 17/17 [00:00<00:00, 184.65it/s]

cost=2.145199, accuracy=0.574300

### 1.1.5 Feedback

Aaaaaand we're done

If you have any suggestions on how we could improve this session, please let us know in the following cell. What did you particularly like or dislike? Did you miss any contents?

### 1.1.6 Additional Tasks

**Task 6a:** Implement the forward function for the SoftMax layer.

**Task 6b:** Implement the estimate function for the CrossEntropy cost.

**Task 6c:** Implement the gradient function for the CrossEntropy cost.

**Task 6d:** Implement the backward function for the SoftMax layer.

**Task 6e:** Test your implementation by setting up a network using the soft max layer and cross entropy cost in combination.

```
In [7]: class SoftMax(Layer):
        def __init__(self):
            self.x_out = None

        def forward(self, x_in):
            # ----- add code for task 6a between comments -----
            # -----
            return self.x_out

        def backward(self, d_out):
            # ----- add code for task 6d between comments -----
            # -----
            return d_in

class CrossEntropy(Cost):
    def __init__(self):
        self.x_in = None
        self.target = None
        self.eps = 1e-12

    def estimate(self, target, x_in):
        # ----- add code for task 6b between comments -----
        # -----
        return cost

    def gradient(self, d_out):
        # ----- add code for task 6c between comments -----
        # -----
        return gradient

In [8]: # define hyperparameters
        input_size = 28*2
        label_size = 10
        batch_size = 600
        learning_rate = 0.00001
        number_epochs = 100

        # ----- add code for task 6e between comments -----
        # -----
```