

# 9\_GAN\_students

July 13, 2019

## 1 Hands-On Machine Learning

### 1.1 Session 9: Generative Adversarial networks

by Christoph Haarburger, Laxmi Gupta

#### 1.1.1 Goals of this Session

In this session you will... \* learn how to implement a vanilla “Goodfellow-like” generative adversarial network \* train a GAN \* learn how to implement a convolutional GAN

We'll be working with the MNIST dataset, which you already know from previous sessions.

```
In [1]: import torch
        import torch.nn as nn
        import torchvision
        import matplotlib.pyplot as plt
        import numpy as np
        import os
os.environ['CUDA_VISIBLE_DEVICES'] = '6'
%matplotlib inline

In [2]: torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False
        torch.manual_seed(0)
        np.random.seed(0)
```

### 1.2 Hyperparameters

```
In [3]: latent_size = 100
        hidden_size = 256
        image_size = 784
        num_epochs = 30
```

**Task:** Set up the necessary transforms that are applied to the input images. You've already done this in the previous sessions.

We'd like to transform our data to Tensors and normalize it with 0.5 mean and 0.5 variance. After defining the transforms, we can initialize the dataset.

```
In [4]: from torchvision import transforms
        from torchvision.datasets import MNIST
        transform = transforms.Compose([
            #transforms.Resize(image_size),
            #transforms.CenterCrop(image_size),
            transforms.ToTensor(),
            transforms.Normalize([0.5], [0.5])
        ])
        dataset = MNIST(root='./data', train=True, download=True, transform=transform)
```

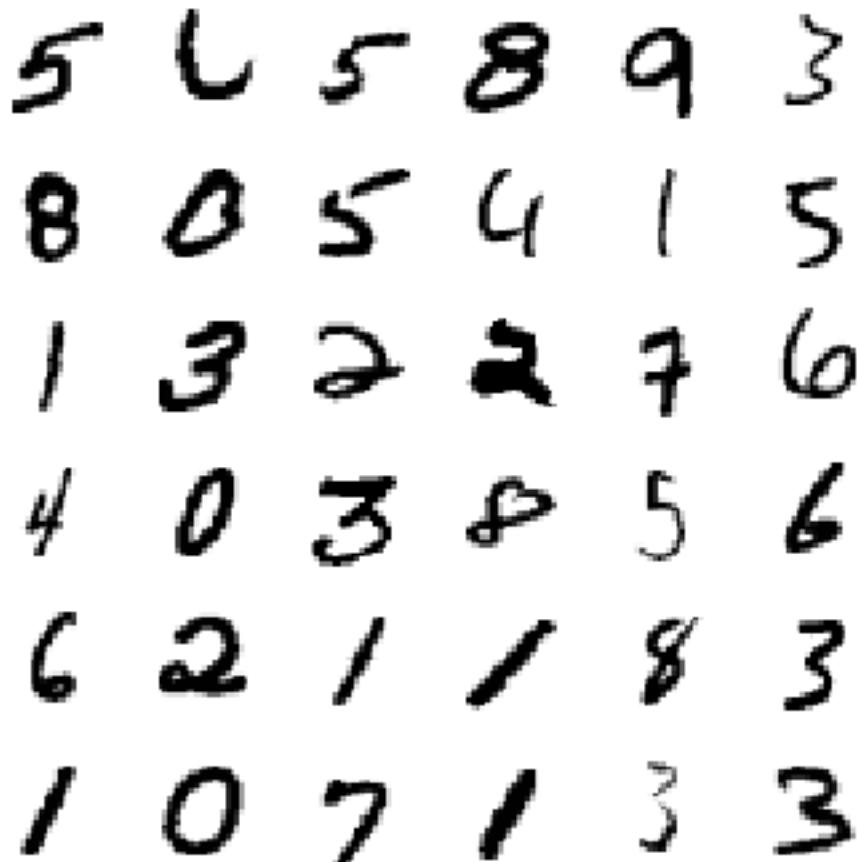
Now we can set up the DataLoader:

```
In [5]: batch_size = 128
        data_loader = torch.utils.data.DataLoader(dataset=dataset,
                                                    batch_size=batch_size,
                                                    shuffle=True, drop_last=True)
```

**Task:** Visualize some samples of a single batch using the plot\_batch() helper function.

```
In [6]: from utils import plot_batch
```

```
for i_batch, sample_batched in enumerate(data_loader):
    plot_batch(sample_batched[0])
    break
```



### 1.3 Build Discriminator

**Task:** Implement the Discriminator as a sequential model consisting of two nn.Linear, nn.ReLU layers and a binary output with nn.Sigmoid activation.

We will feed the images into the network as flattened arrays.

```
In [7]: discriminator = nn.Sequential(  
    nn.Linear(image_size, hidden_size),  
    nn.ReLU(),  
    nn.Linear(hidden_size, 10),  
    nn.ReLU(),  
    nn.Linear(10, 1),  
    nn.Sigmoid()  
)
```

### 1.4 Build Generator

**Task:** Implement the Generator as a sequential model of two nn.Linear layers with nn.ReLU() activation and a final nn.Linear layer with nn.Tanh activation that maps from a random vector to the original flattened image dimension.

```
In [8]: generator = nn.Sequential(  
    nn.Linear(latent_size, hidden_size),  
    nn.LeakyReLU(),  
    nn.Linear(hidden_size, hidden_size),  
    nn.LeakyReLU(),  
    nn.Linear(hidden_size, image_size),  
    nn.Tanh()  
)
```

#### 1.4.1 General Setup for Training

By the following line we can dynamically determine whether we want to run our model and optimization on CPU or GPU

```
In [9]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

**Task:** Send both generator and discriminator to the device. You can do that simply by calling {mymodule}.to(device)

```
In [10]: netG = generator.to(device)  
netD = discriminator.to(device)
```

**Task:** Now we can initialize the optimizers for both generator and discriminator. We'll use Adam with a learning rate of 0.0005 here.

```
In [11]: from torch.optim import Adam
d_optimizer = Adam(netD.parameters(), lr=0.0005)
g_optimizer = Adam(netG.parameters(), lr=0.0005)
```

Now we only need to initialize the cross entropy loss before getting started with the training loop

Compute BCE\_Loss using real images where  $\text{BCE\_Loss} = -y * \log(D(x)) - (1-y) * \log(1 - D(x))$ . The second term of the loss is always zero since `real_labels == 1`.

```
In [12]: criterion = nn.BCELoss()
```

## 1.4.2 Training Loop

Finally, we can write the training loop and look at actual results. Since this is quite a bit of code, please perform the implementation in the following steps:

**Task:** Train your GAN and watch your fake images get better and better. It might take up to 20 Minutes to produce good results.

```
In [13]: total_step = len(data_loader)
for epoch in range(num_epochs):
    for i, (images, _) in enumerate(data_loader):
        images = images.reshape((images.size(0), -1)).to(device)

        # Create the labels which are later used as input for the BCE loss
        real_labels = torch.ones(batch_size, 1).to(device)
        fake_labels = torch.zeros(batch_size, 1).to(device)

        # ===== #
        #           Train the discriminator                      #
        # ===== #

        outputs = discriminator(images)
        d_loss_real = criterion(outputs, real_labels)
        real_score = outputs

        # Compute BCELoss using fake images
        z = torch.randn(batch_size, latent_size).to(device)
        fake_images = generator(z)
        outputs = discriminator(fake_images)
        d_loss_fake = criterion(outputs, fake_labels)
        fake_score = outputs

        # Backprop and optimize
        d_loss = d_loss_real + d_loss_fake
        d_optimizer.zero_grad()
        g_optimizer.zero_grad()
        d_loss.backward()
        d_optimizer.step()
```

```

# ===== #
#           Train the generator                      #
# ===== #

# Compute loss with fake images
z = torch.randn(batch_size, latent_size).to(device)
fake_images = generator(z)
outputs = discriminator(fake_images)

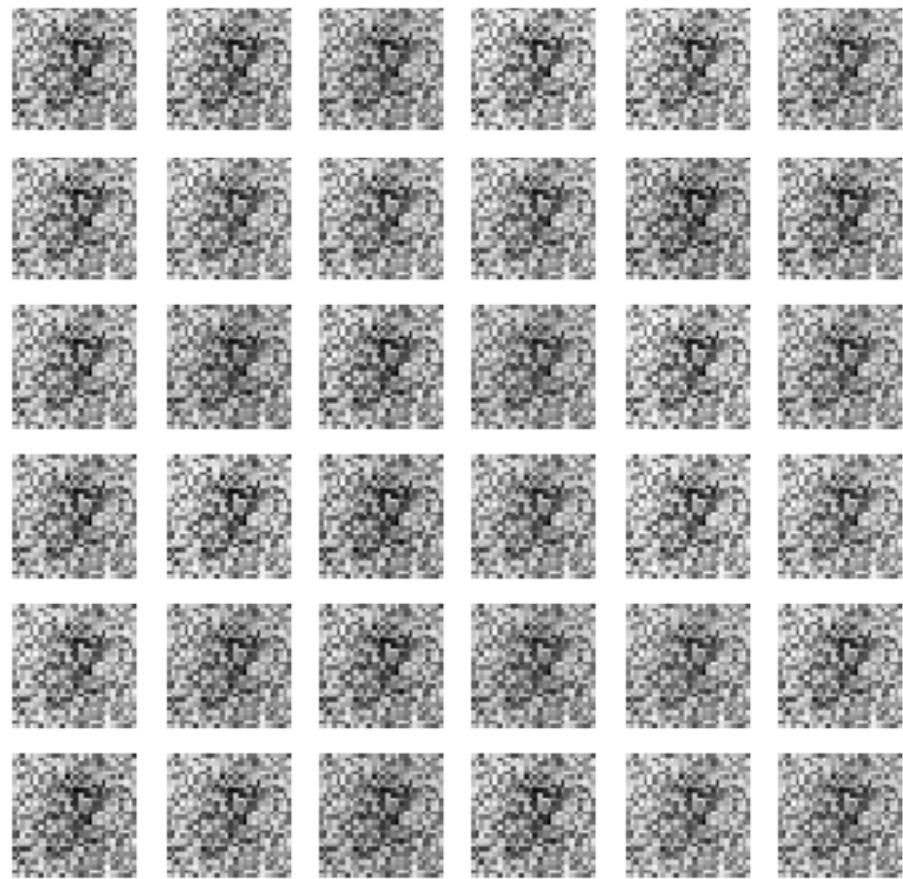
# We train G to maximize log(D(G(z)) instead of minimizing log(1-D(G(z)))
g_loss = criterion(outputs, real_labels)

# Backprop and optimize
d_optimizer.zero_grad()
g_optimizer.zero_grad()
g_loss.backward()
g_optimizer.step()

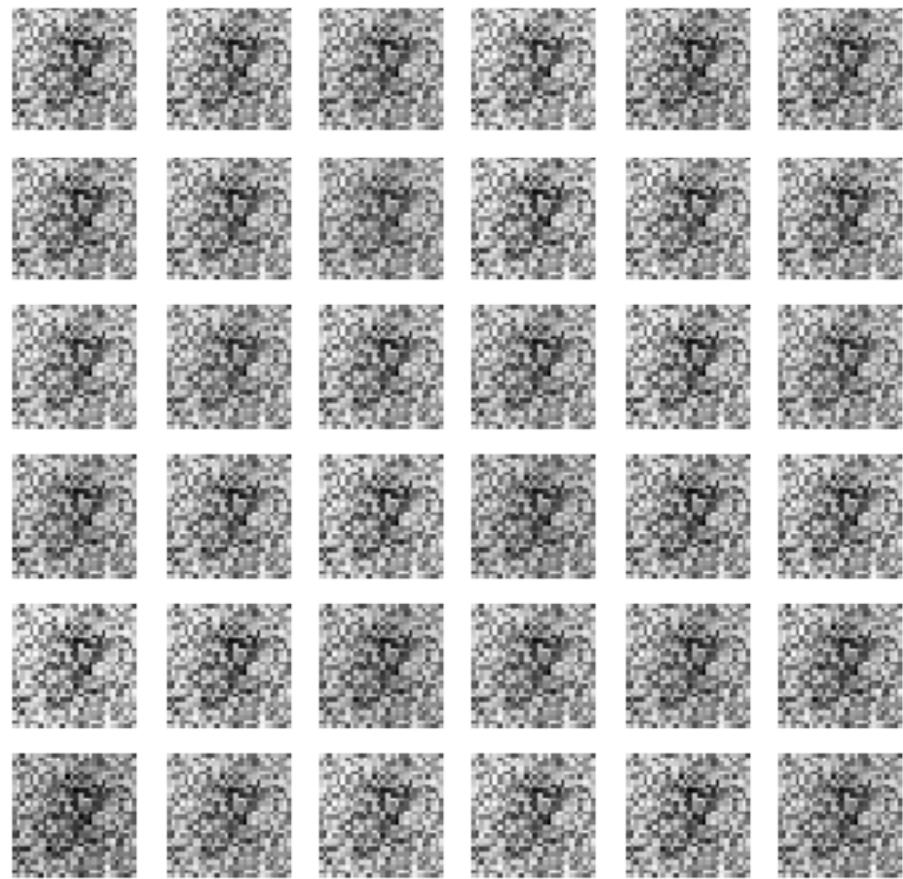
if (i+1) % 200 == 0:
    print('Epoch [{}/{}], Step [{}/{}], d_loss: {:.4f}, g_loss: {:.4f}, D(x): {:.4f}, D(G(z)): {:.4f}'.format(epoch, num_epochs, i+1, total_step, d_loss.item(), g_loss.item(),
                                                       real_score.mean().item(), fake_score.mean().item()))
fake_images = fake_images.reshape(fake_images.size(0), 1, 28, 28).cpu().detach().numpy()
plot_batch(fake_images)

```

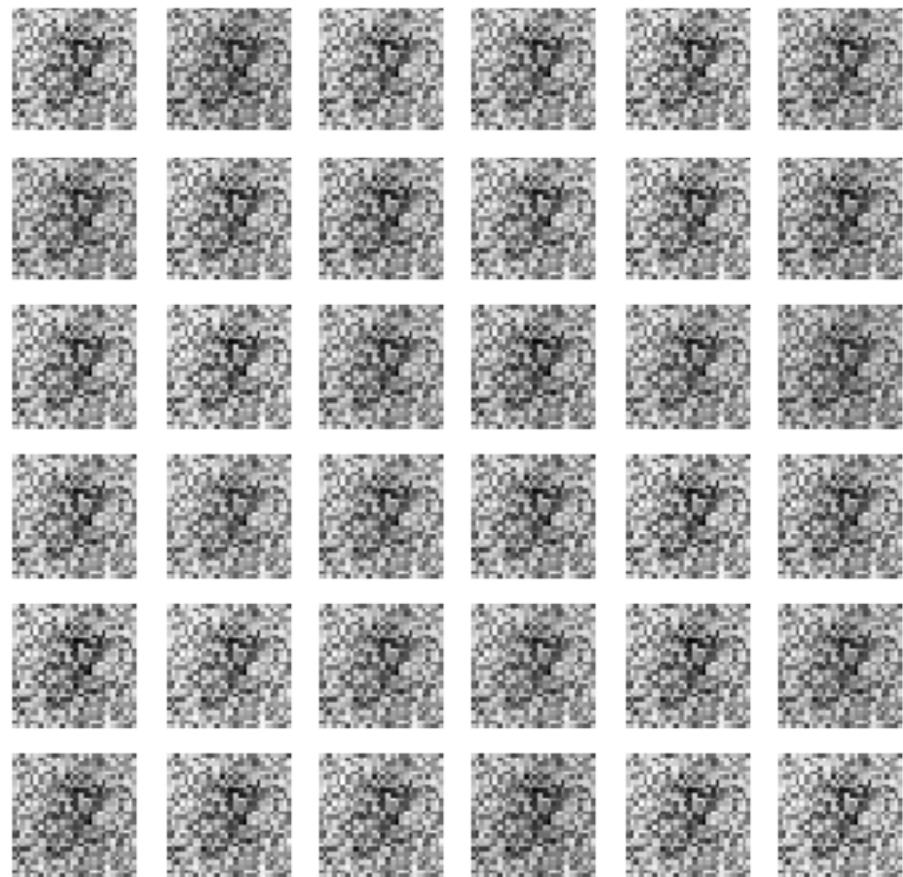
Epoch [0/30], Step [200/468], d\_loss: 1.3867, g\_loss: 0.7146, D(x): 0.49, D(G(z)): 0.49  
Epoch [0/30], Step [400/468], d\_loss: 1.3866, g\_loss: 0.7102, D(x): 0.49, D(G(z)): 0.49



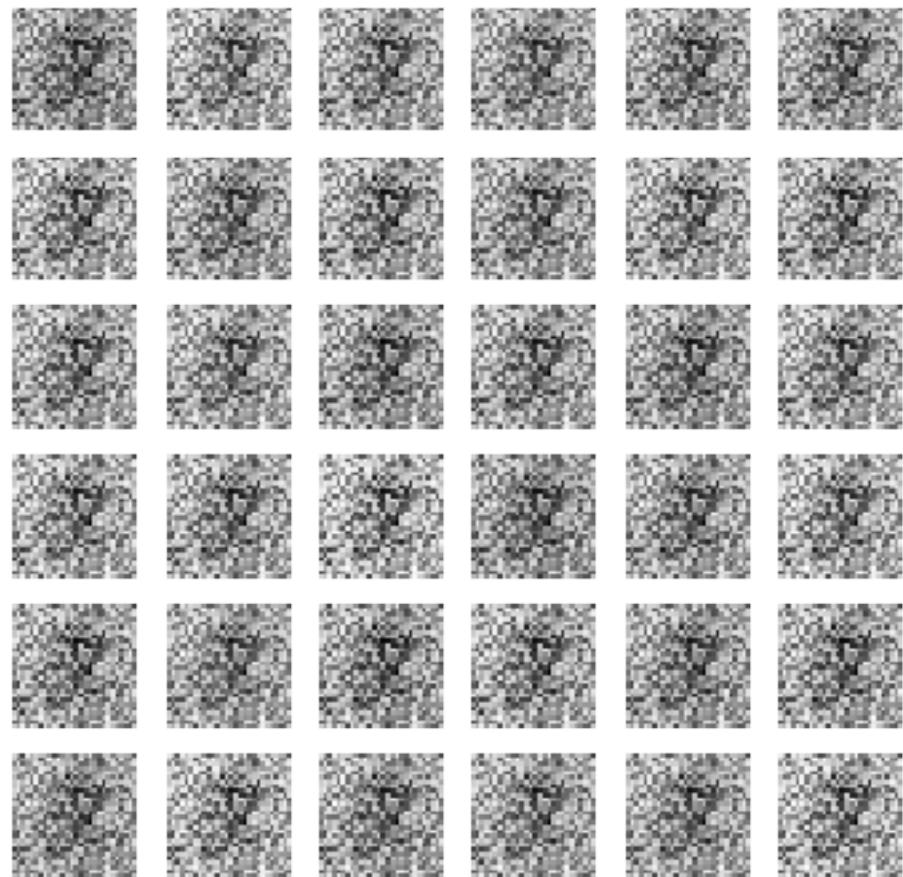
Epoch [1/30], Step [200/468], d\_loss: 1.3864, g\_loss: 0.7045, D(x): 0.49, D(G(z)): 0.49  
Epoch [1/30], Step [400/468], d\_loss: 1.3864, g\_loss: 0.7008, D(x): 0.50, D(G(z)): 0.50



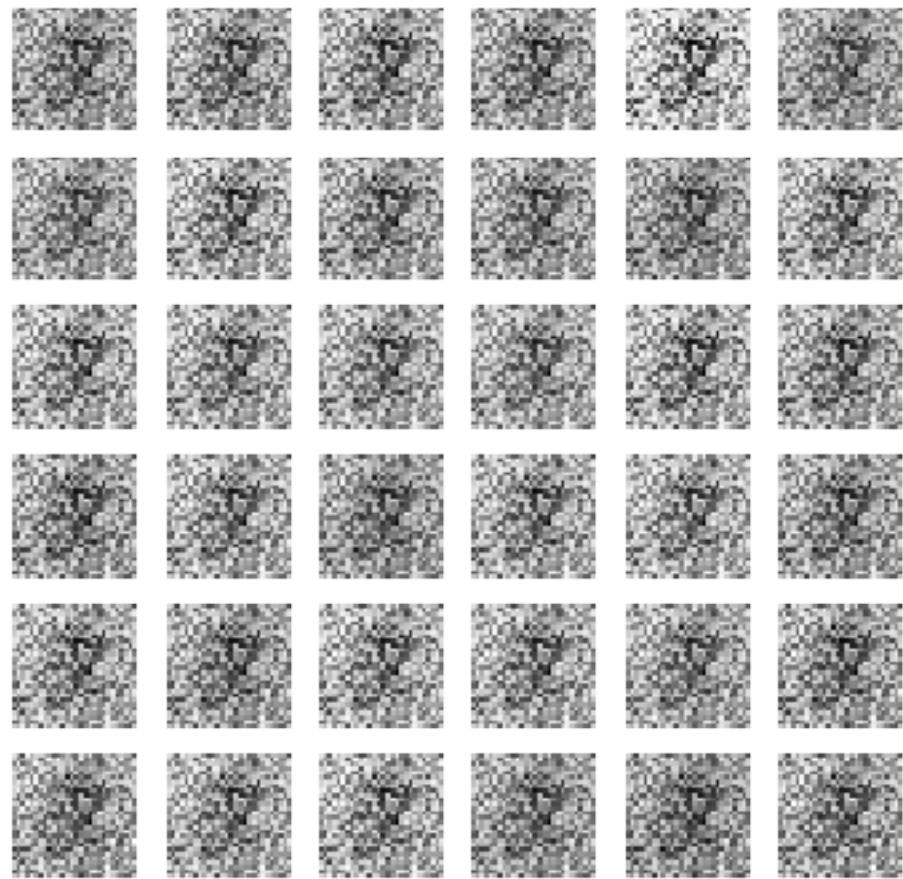
Epoch [2/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6972, D(x): 0.50, D(G(z)): 0.50  
Epoch [2/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6954, D(x): 0.50, D(G(z)): 0.50



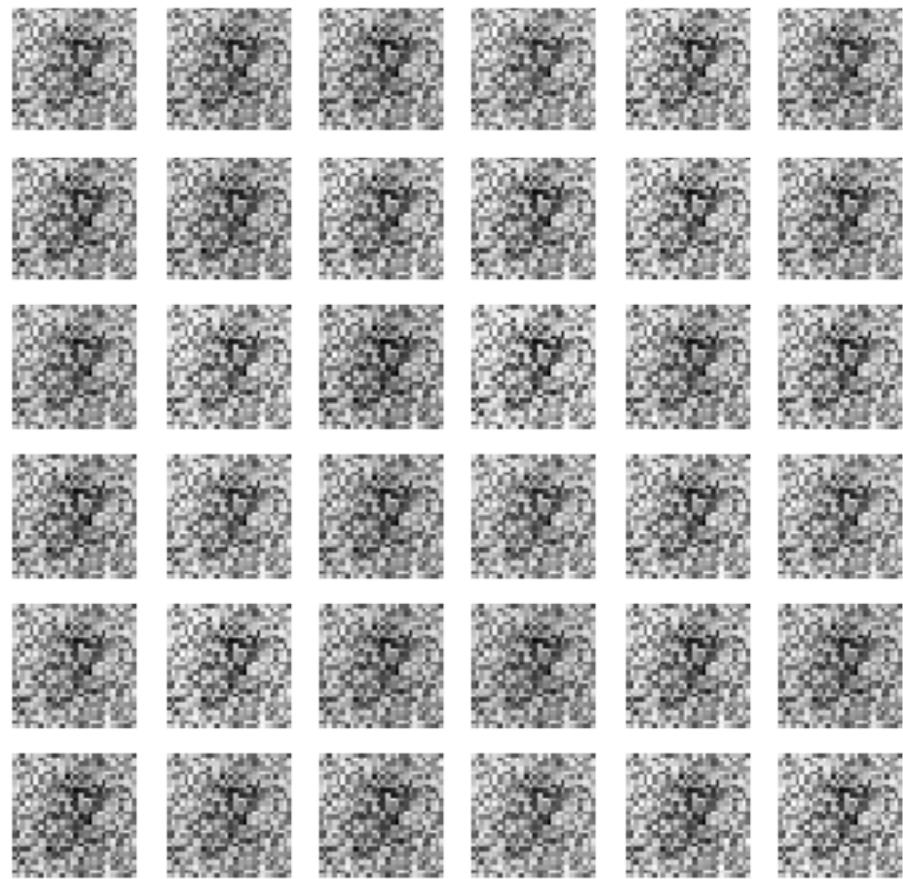
Epoch [3/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6941, D(x): 0.50, D(G(z)): 0.50  
Epoch [3/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6936, D(x): 0.50, D(G(z)): 0.50



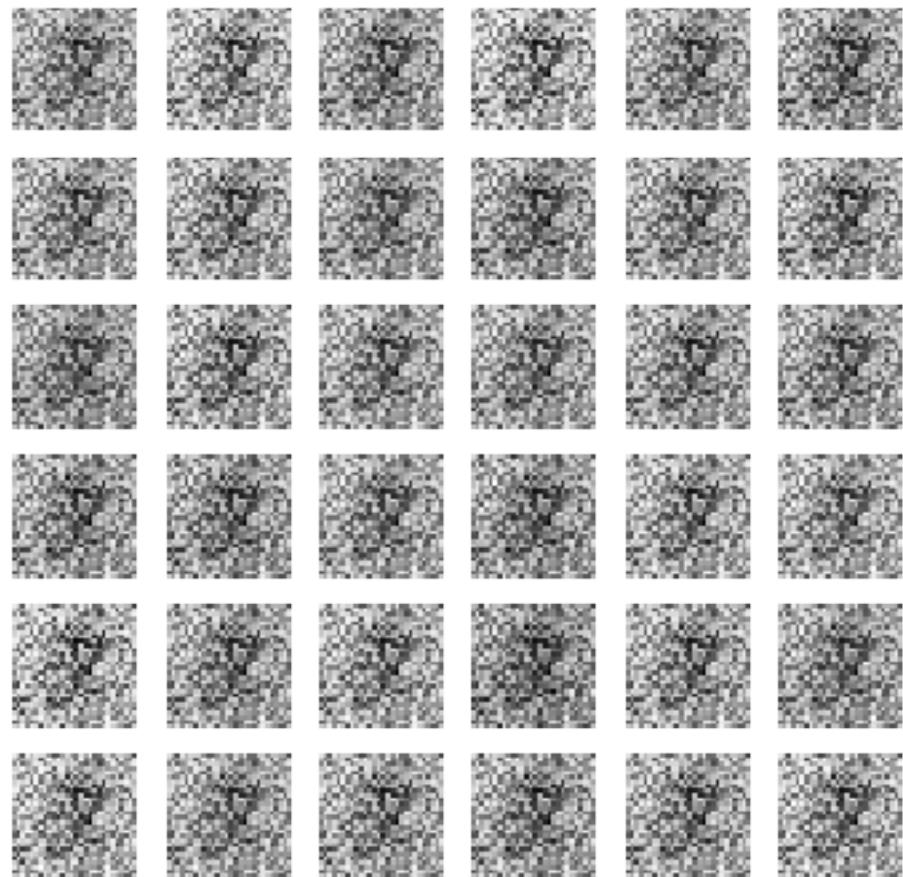
Epoch [4/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6933, D(x): 0.50, D(G(z)): 0.50  
Epoch [4/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6932, D(x): 0.50, D(G(z)): 0.50



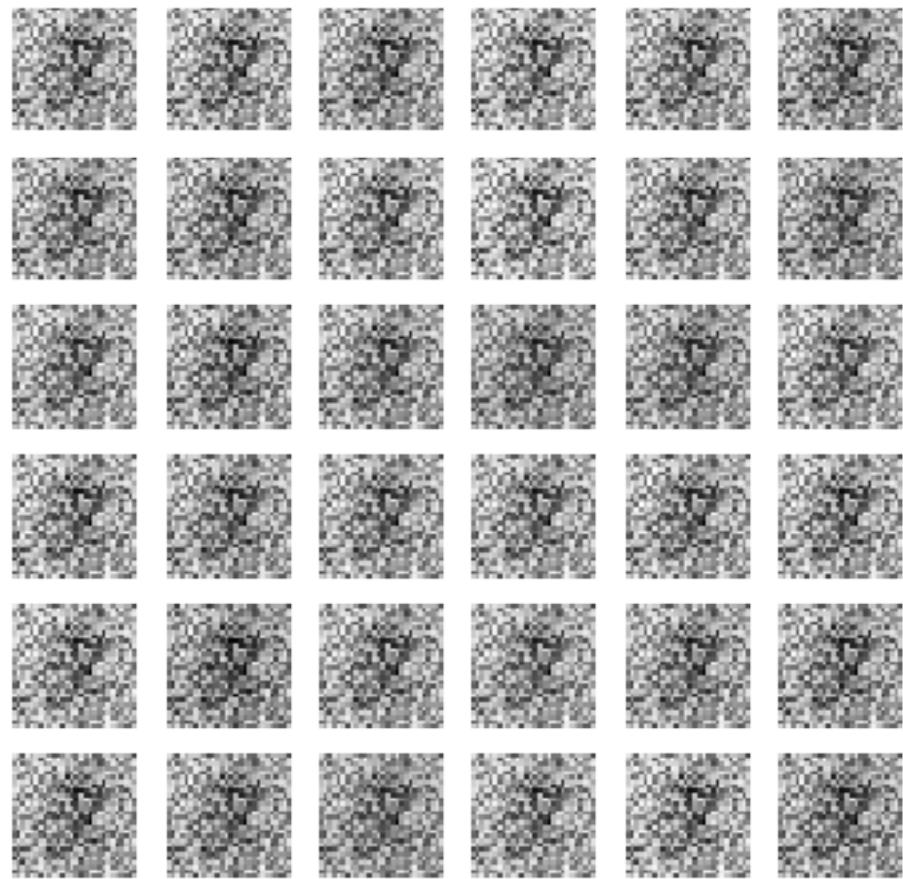
Epoch [5/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6932, D(x): 0.50, D(G(z)): 0.50  
Epoch [5/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



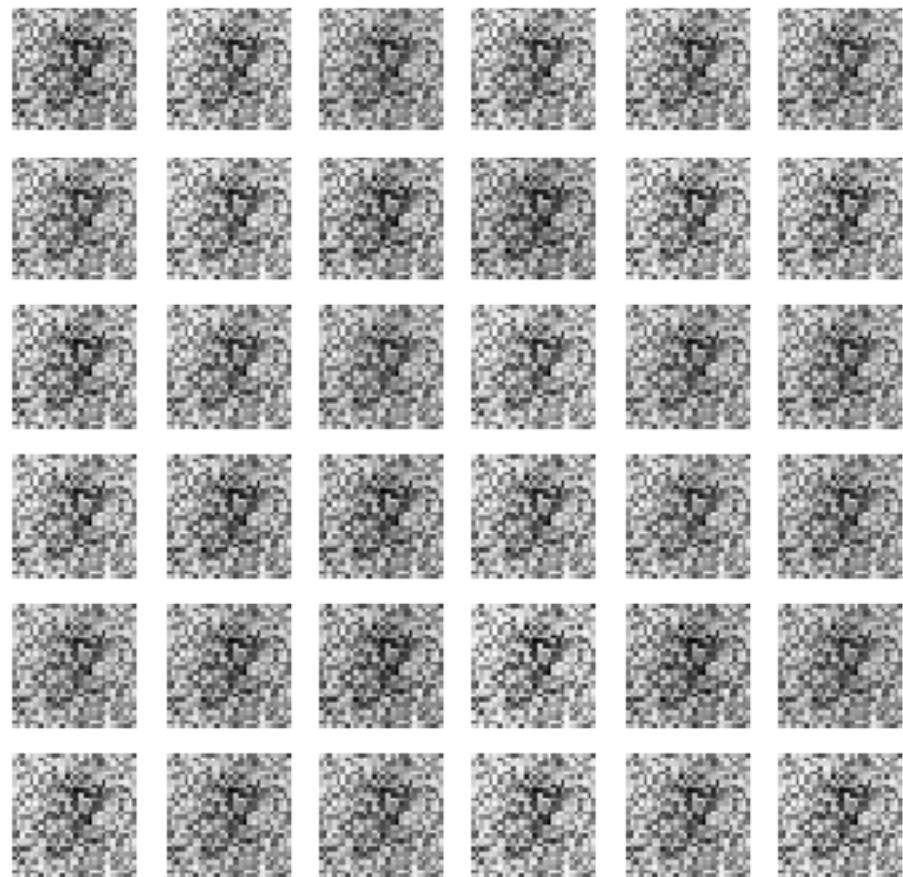
Epoch [6/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [6/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



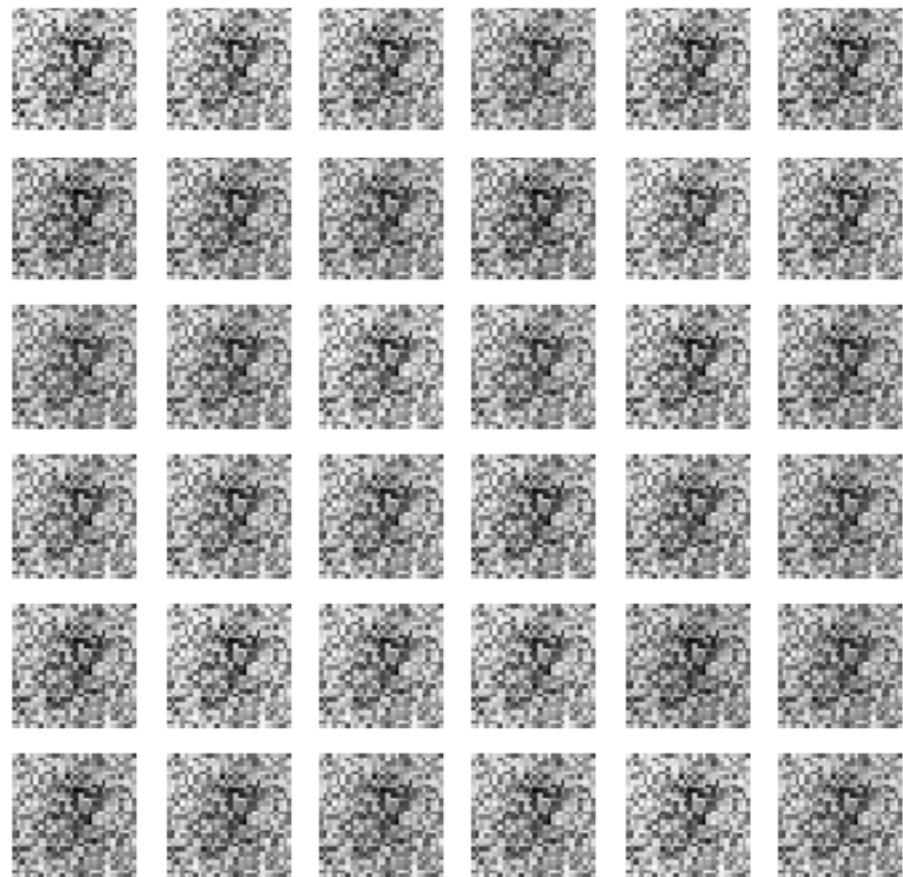
Epoch [7/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [7/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



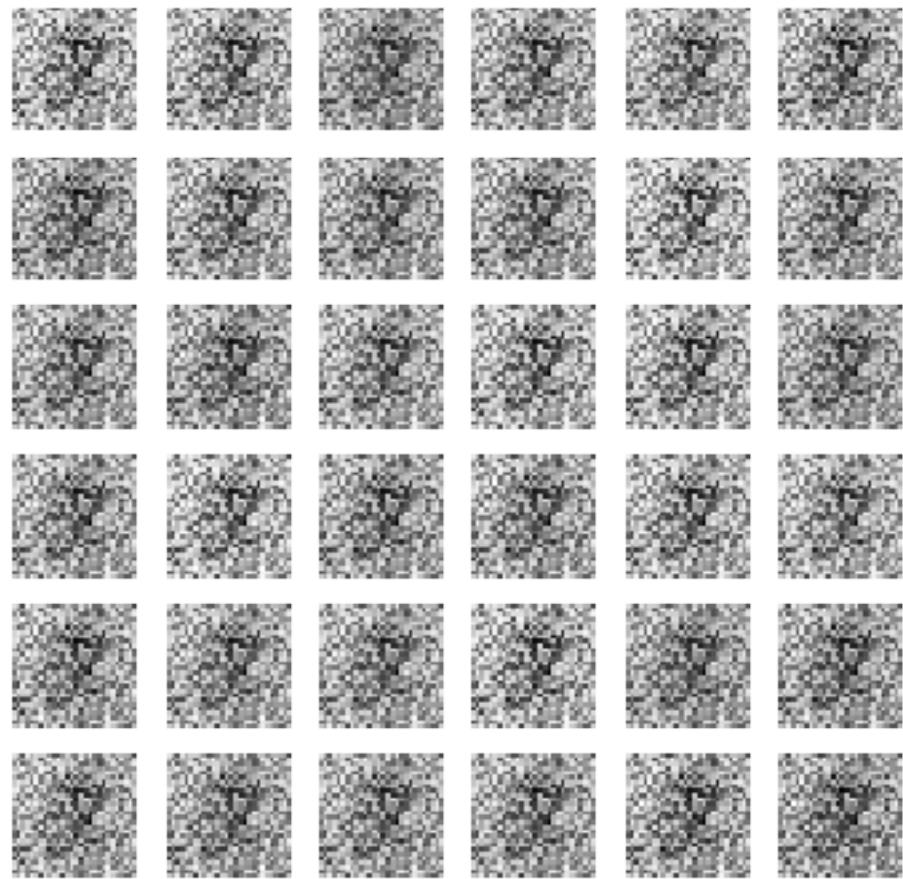
Epoch [8/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [8/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



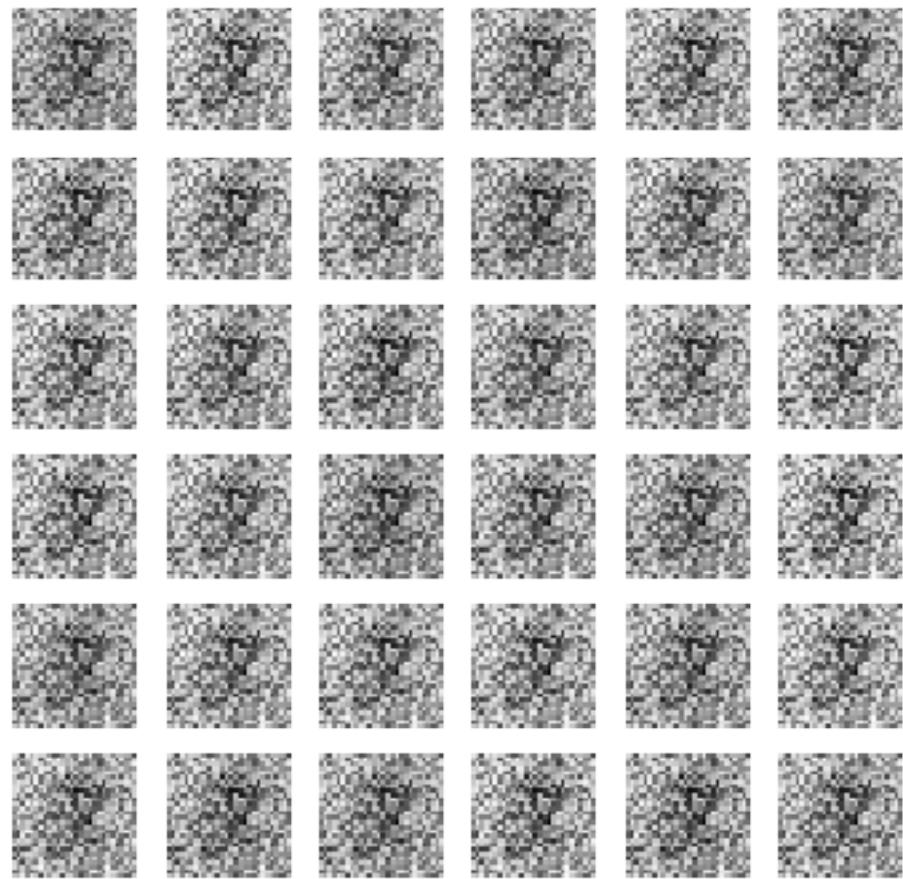
Epoch [9/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [9/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



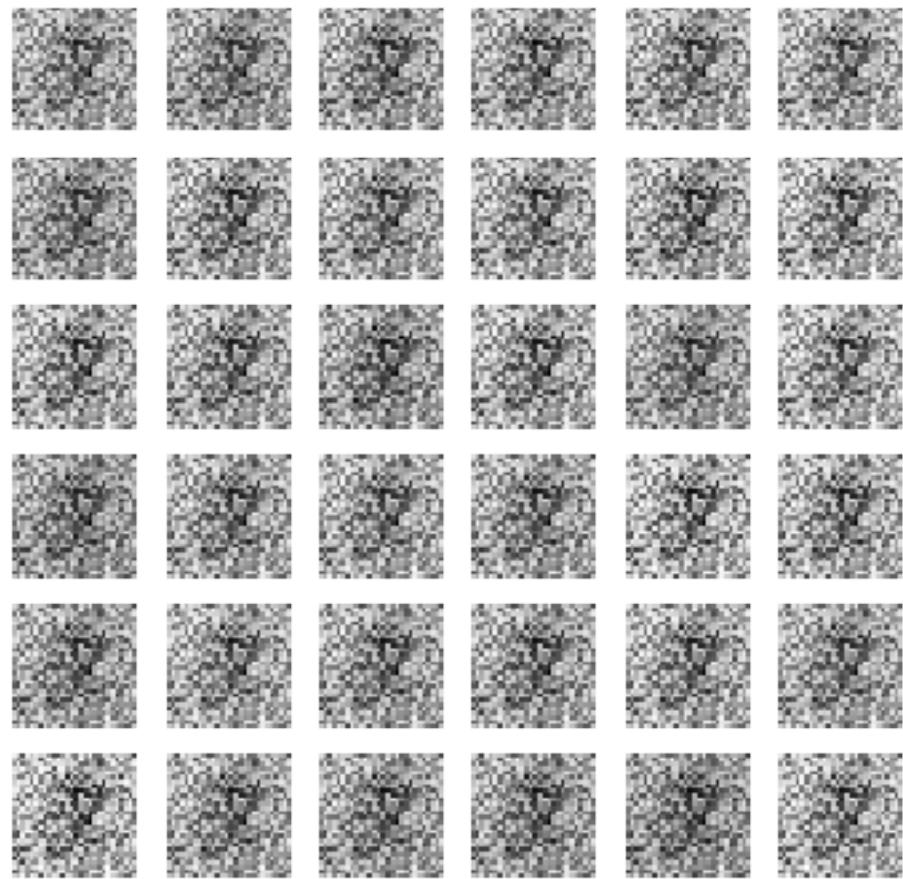
```
Epoch [10/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [10/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



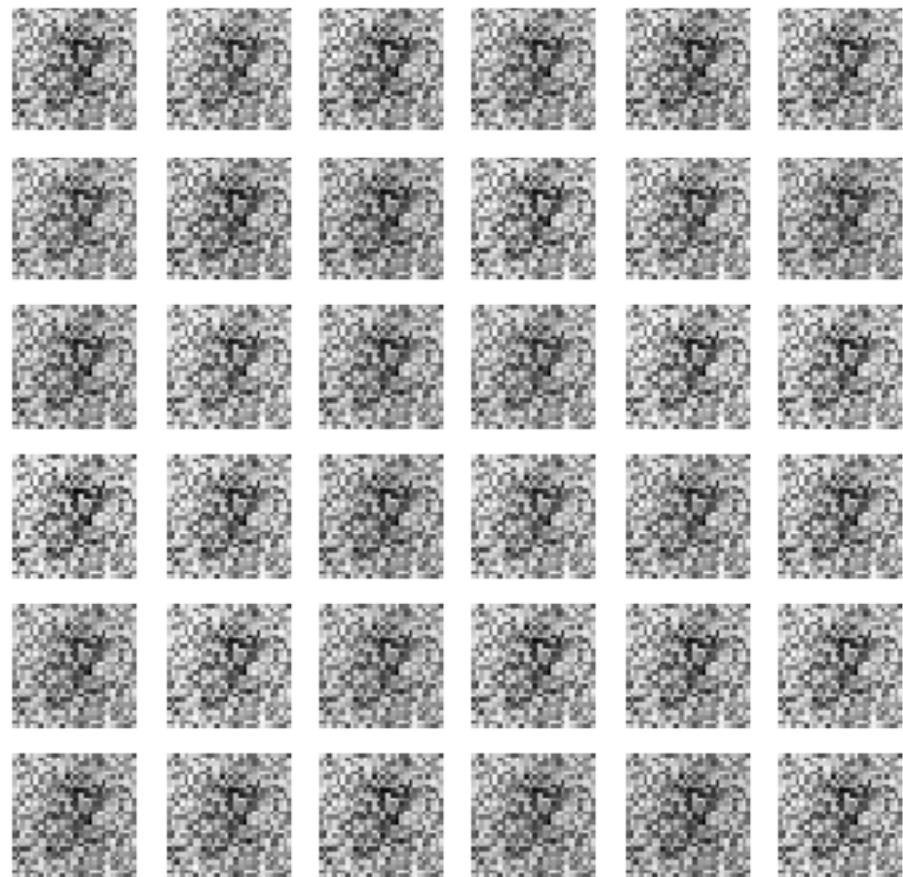
```
Epoch [11/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [11/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



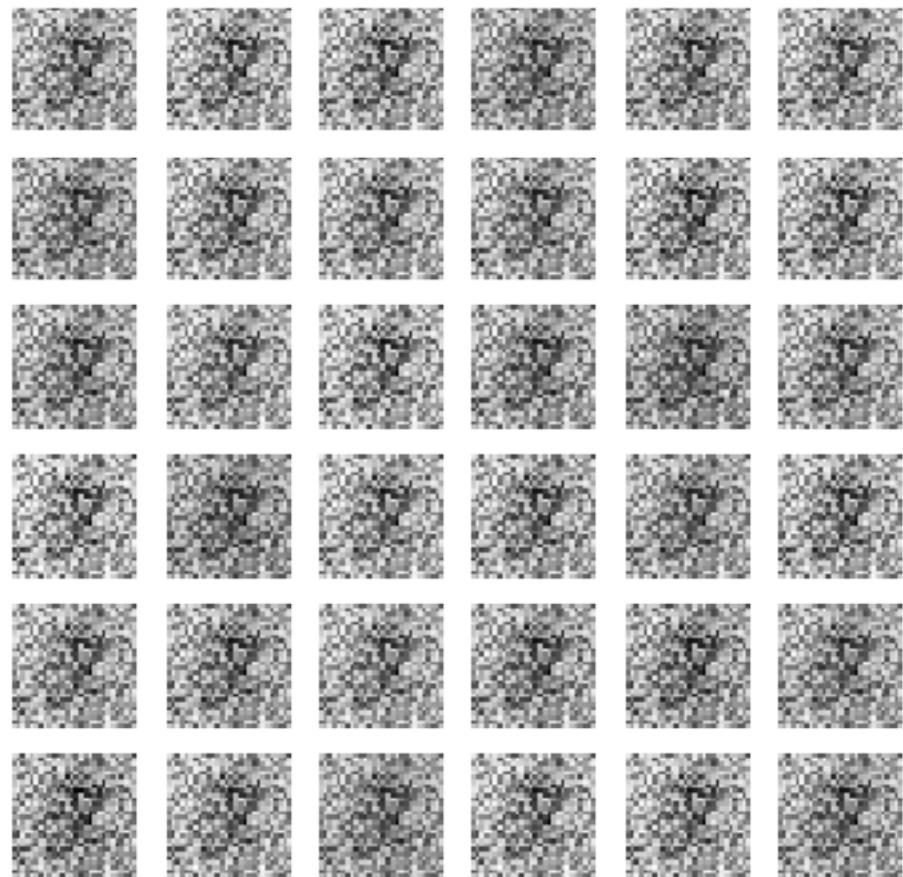
Epoch [12/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [12/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



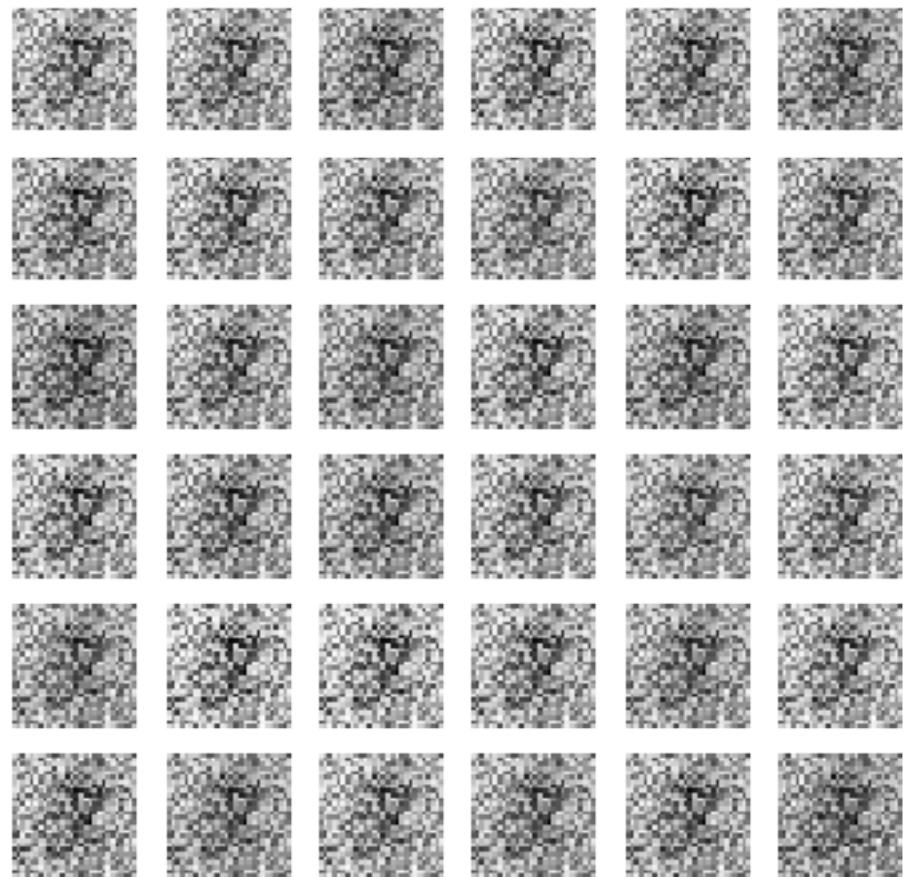
```
Epoch [13/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [13/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



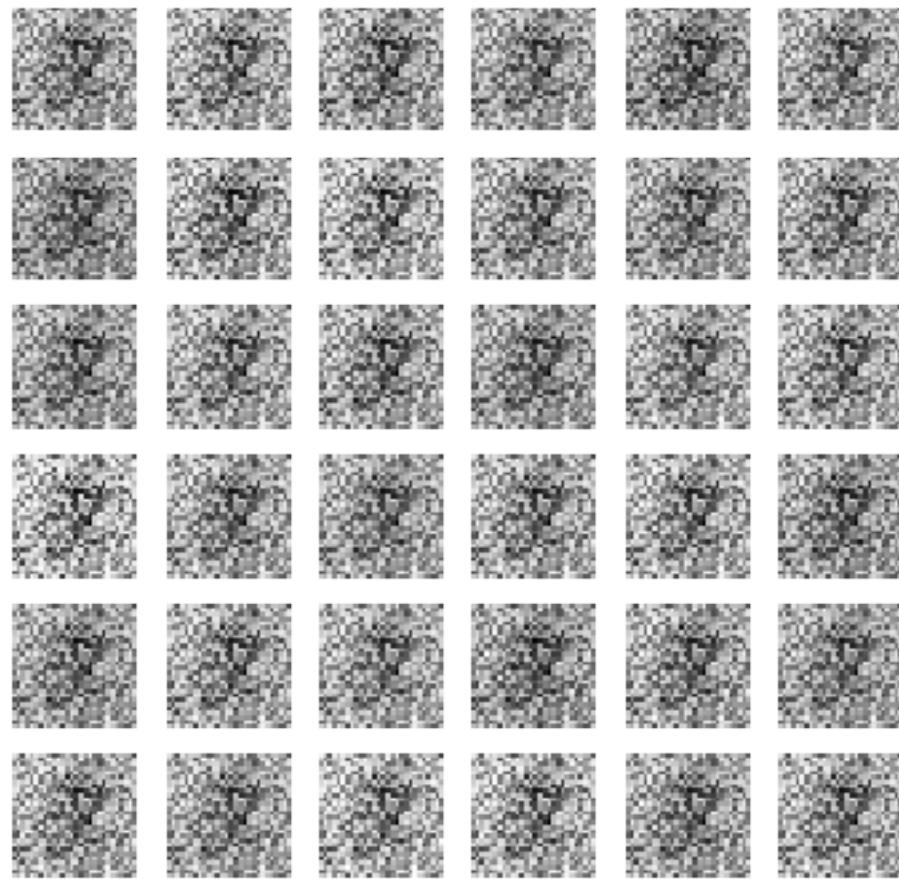
```
Epoch [14/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [14/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



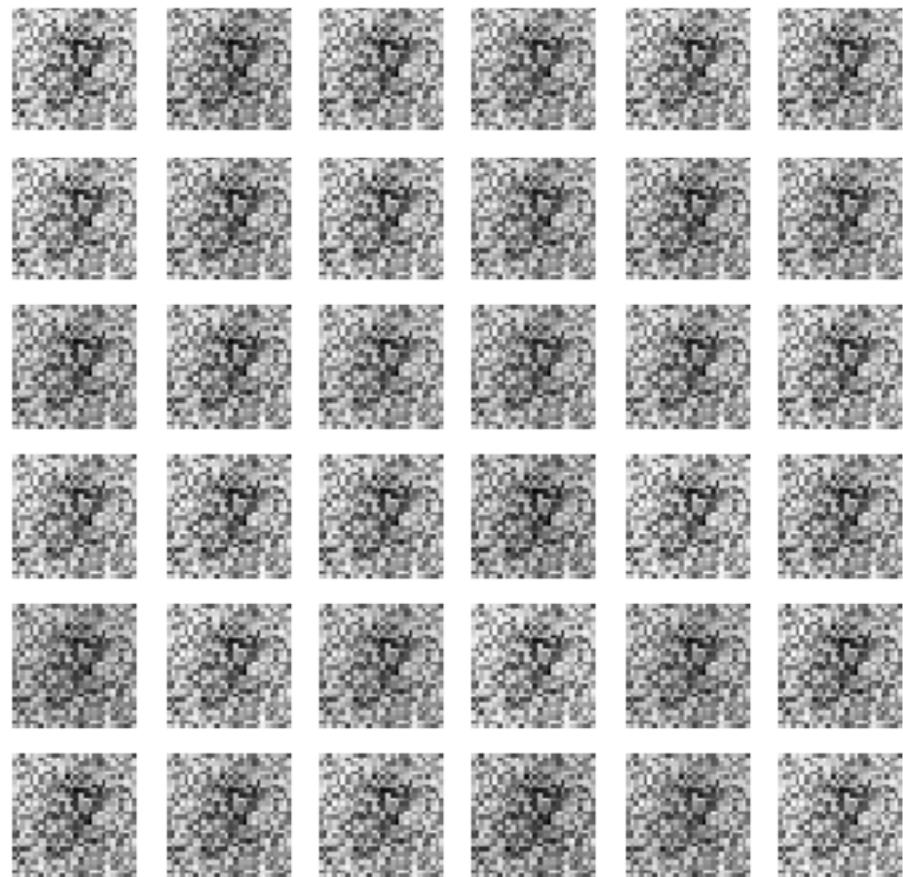
Epoch [15/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [15/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



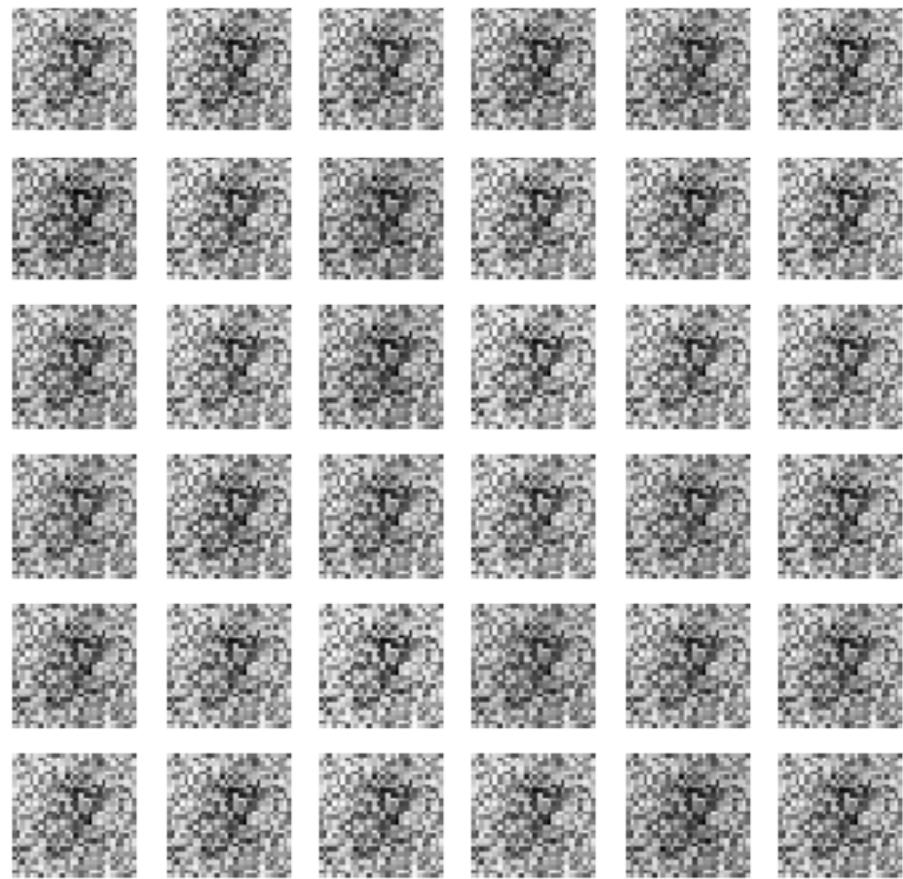
Epoch [16/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [16/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



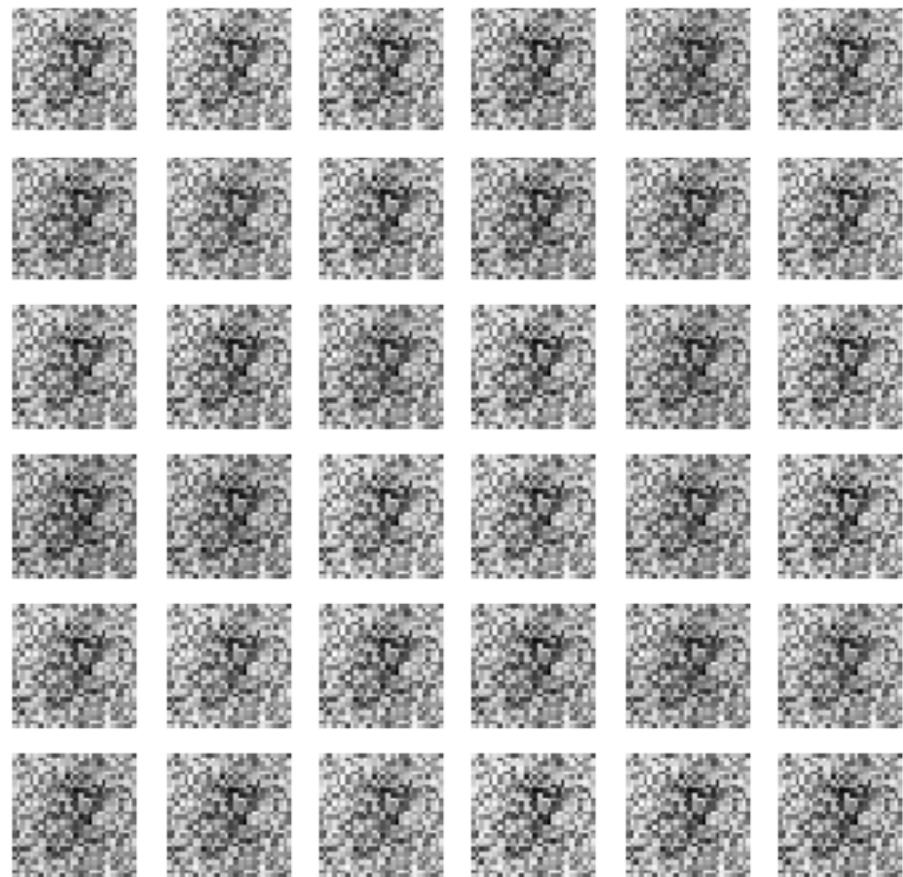
Epoch [17/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [17/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



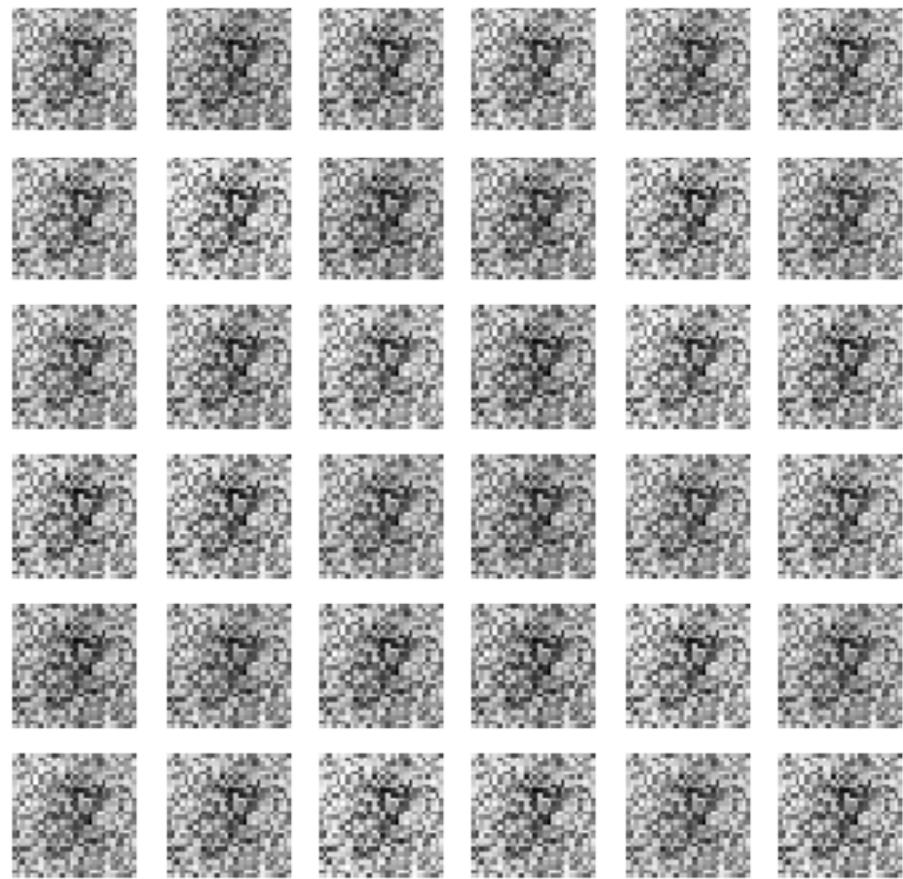
```
Epoch [18/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [18/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



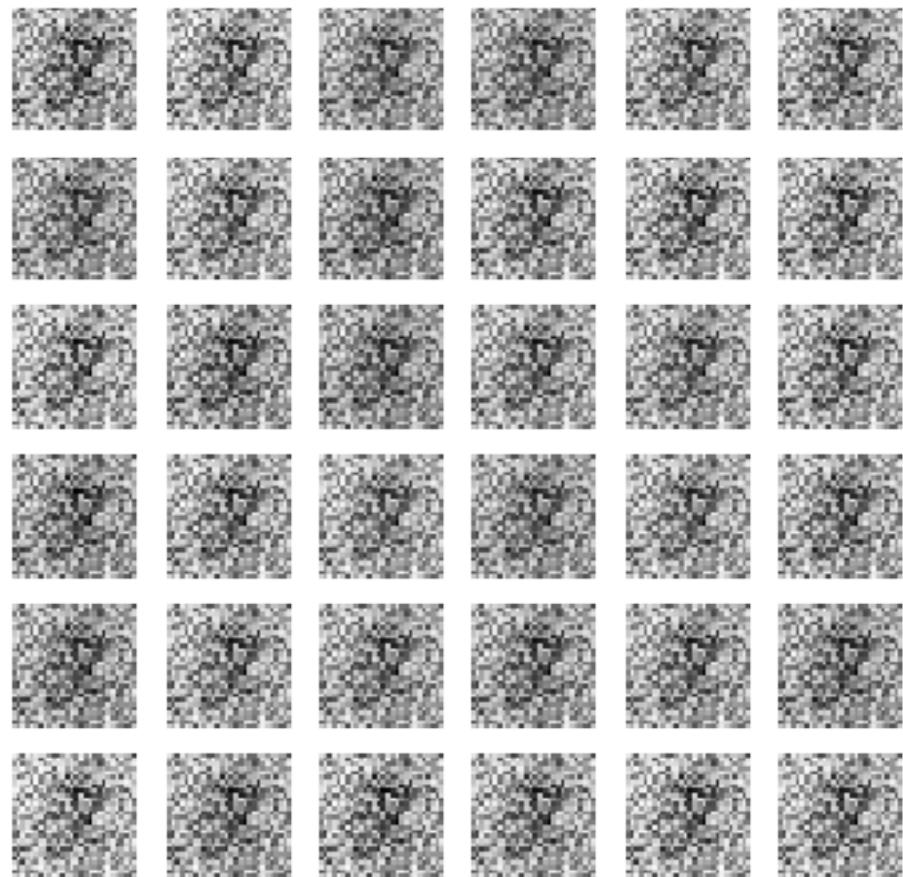
Epoch [19/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [19/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



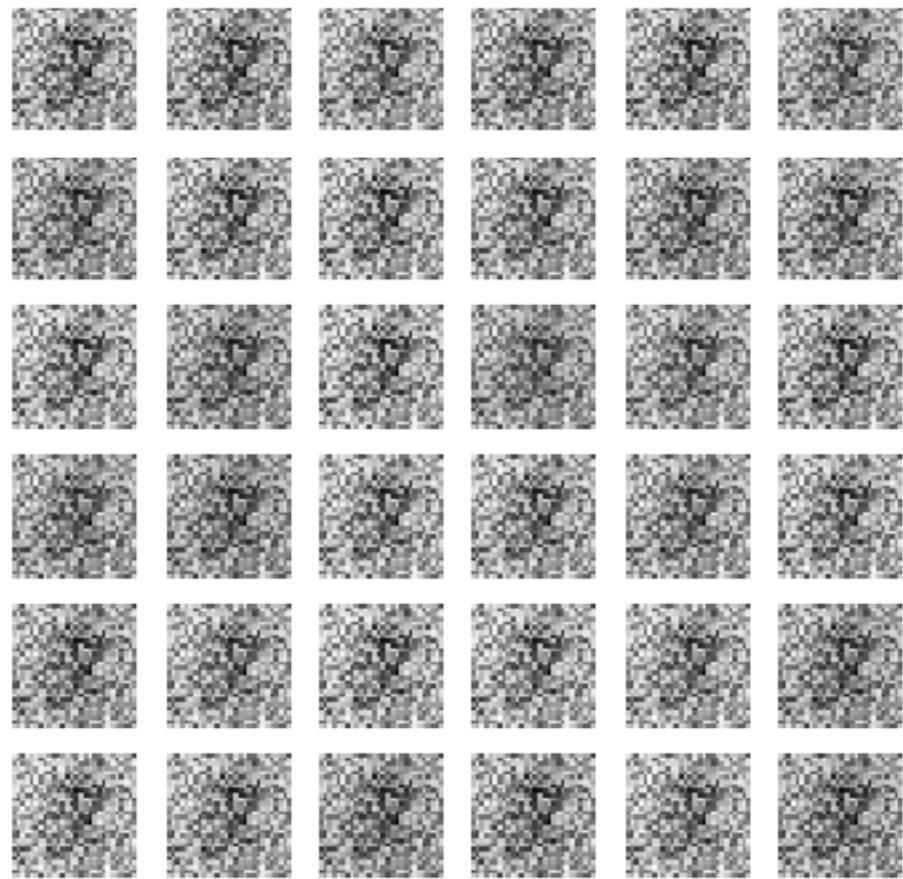
```
Epoch [20/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [20/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



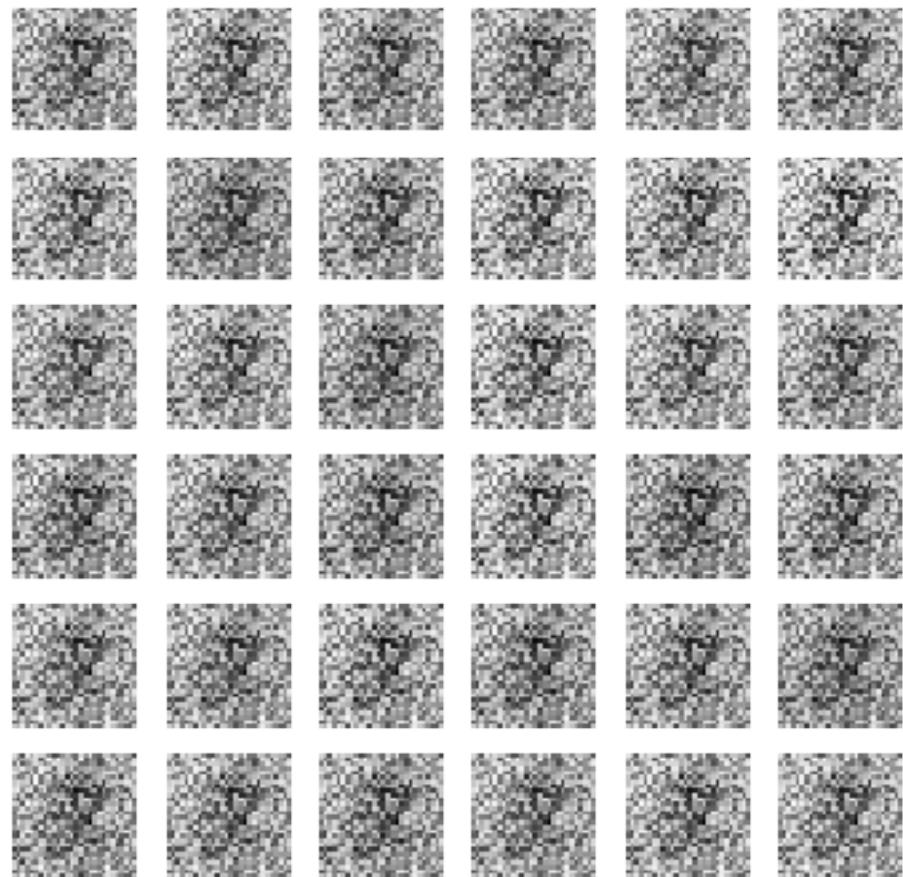
```
Epoch [21/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [21/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



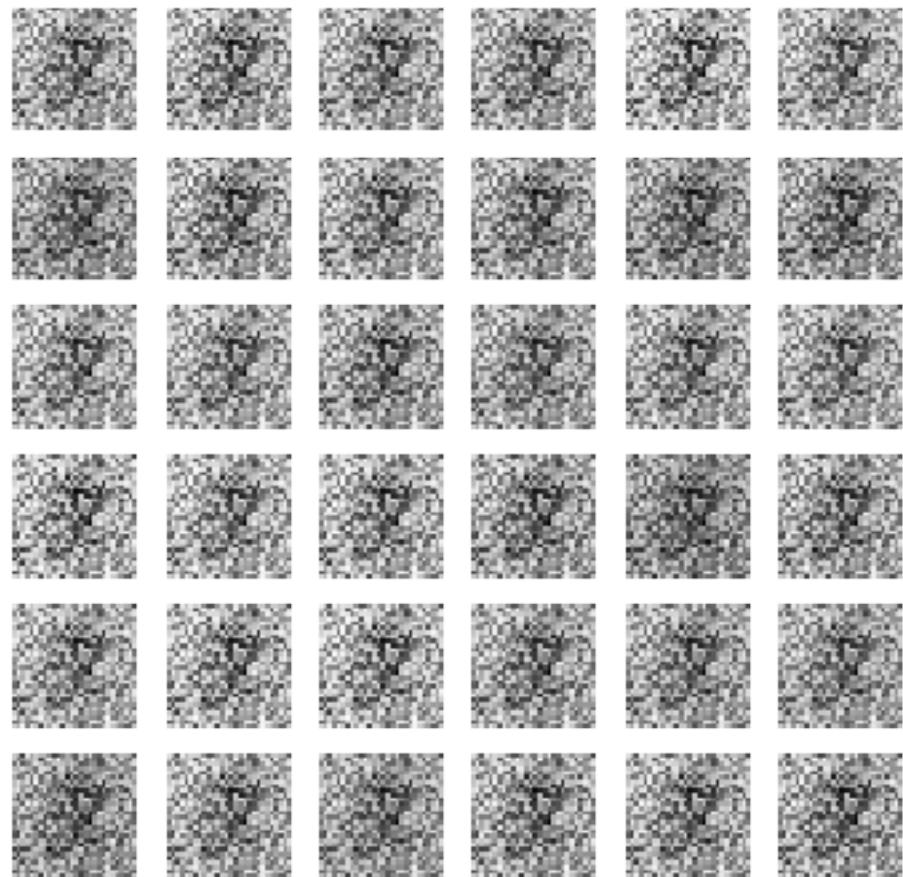
```
Epoch [22/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [22/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



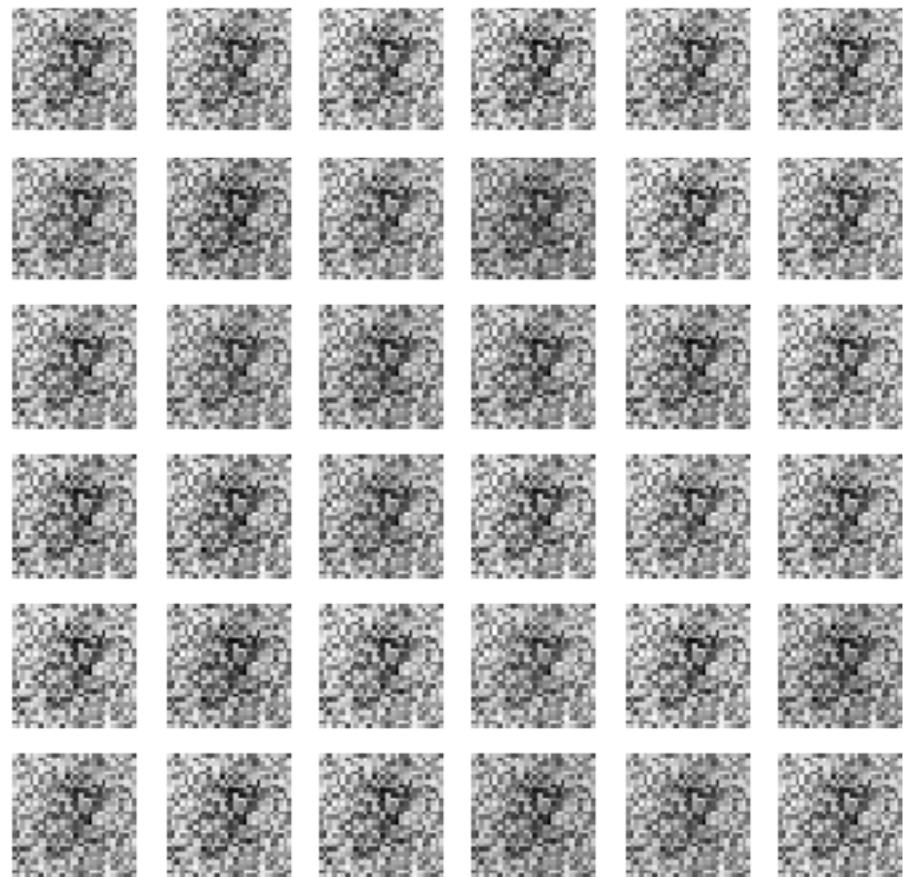
Epoch [23/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [23/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



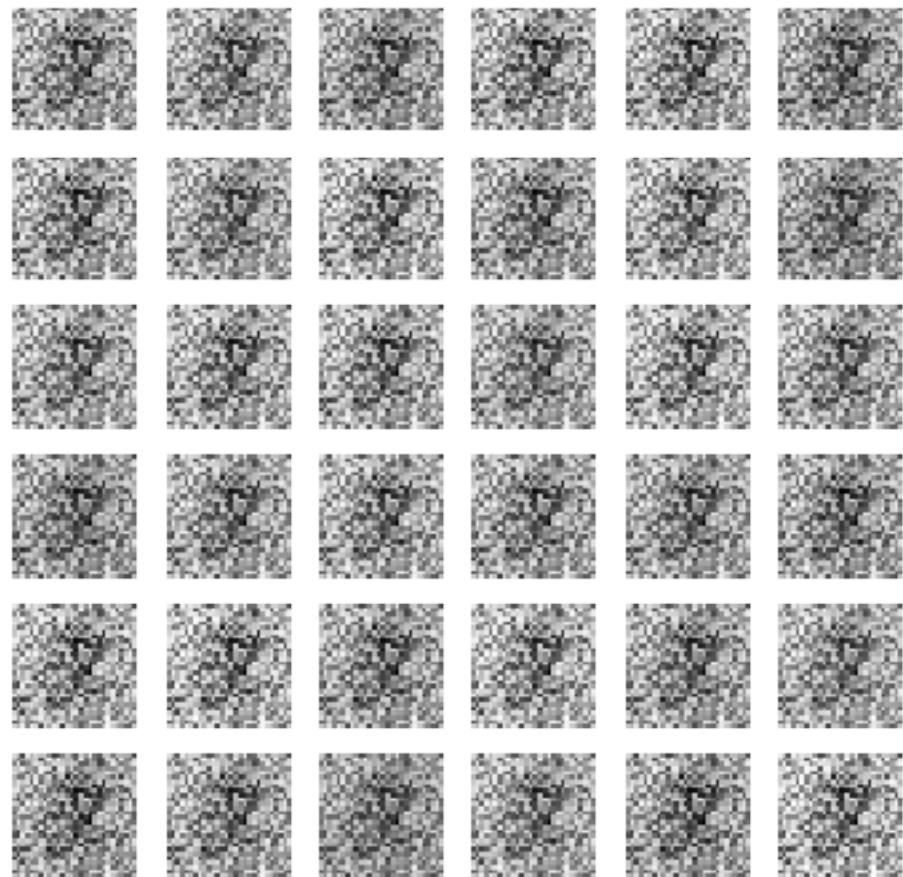
```
Epoch [24/30], Step [200/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
Epoch [24/30], Step [400/468], d_loss: 1.3863, g_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50
```



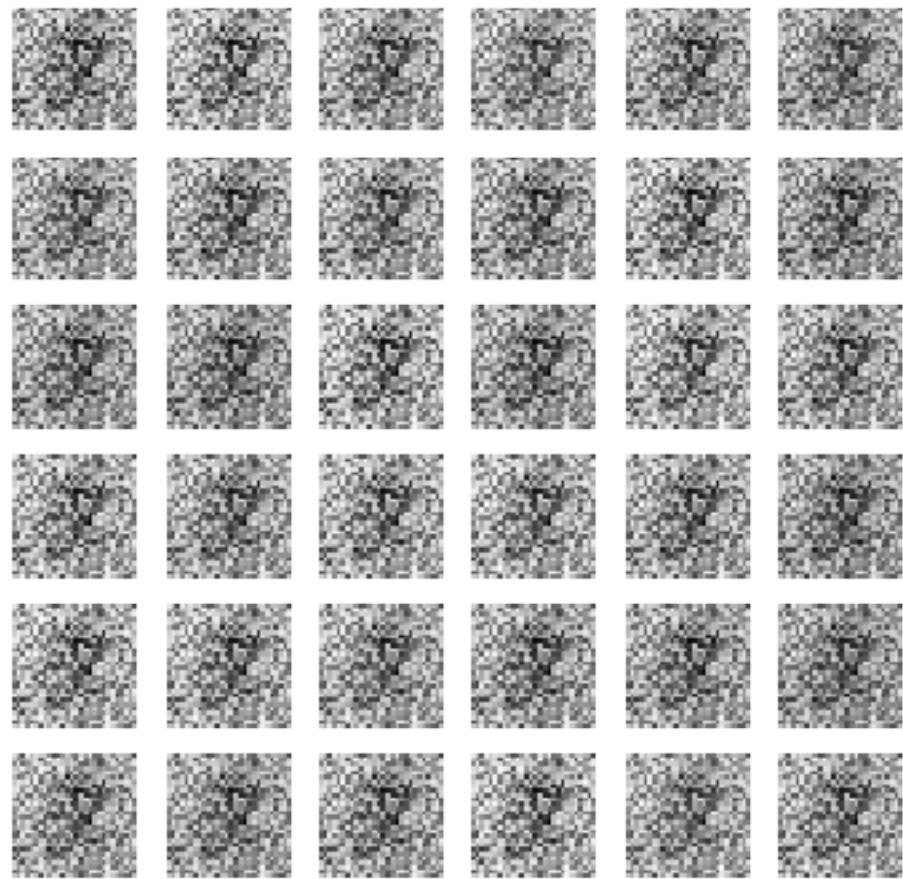
Epoch [25/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [25/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



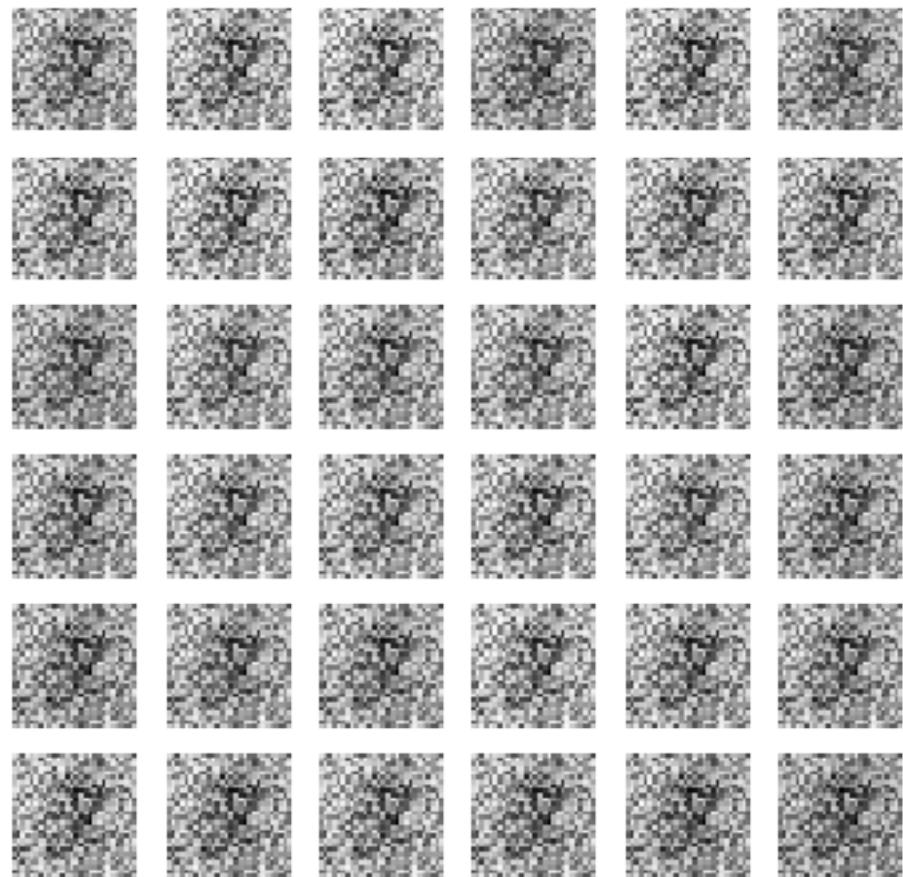
Epoch [26/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [26/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



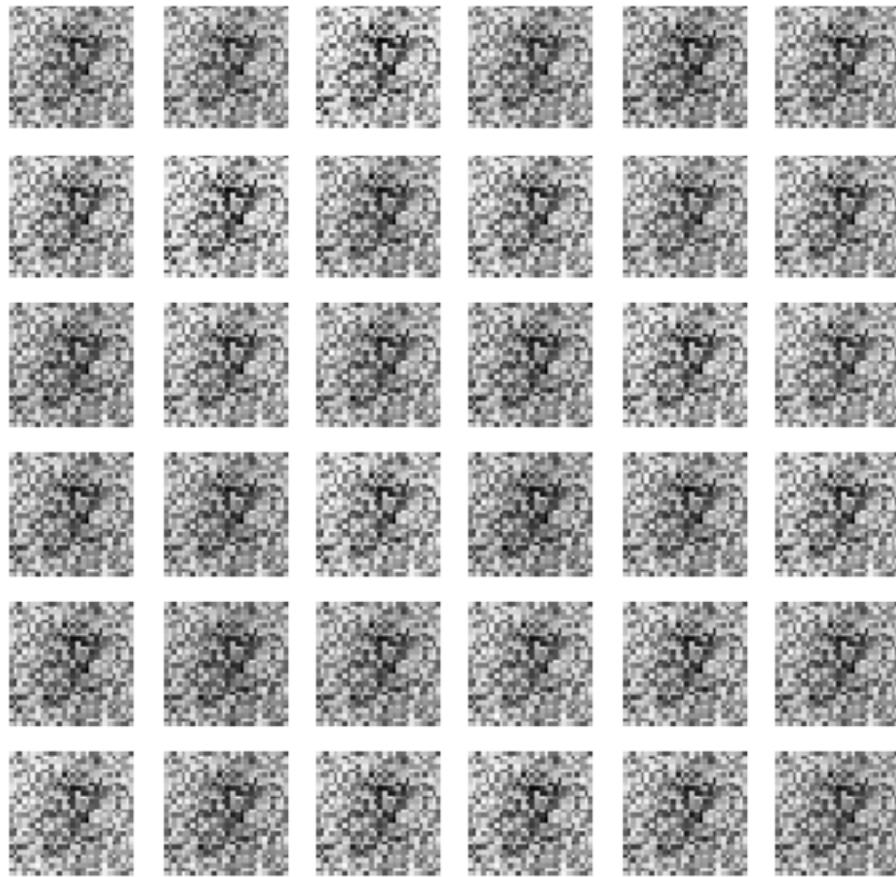
Epoch [27/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [27/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



Epoch [28/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [28/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



Epoch [29/30], Step [200/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50  
Epoch [29/30], Step [400/468], d\_loss: 1.3863, g\_loss: 0.6931, D(x): 0.50, D(G(z)): 0.50



Vanilla GANs as implemented above are difficult to train because of the high number of the many feedforward connections. DCGAN aims to overcome this by using Convolutional instead of Fully-connected layers.

In the next cells we will implement the same Training loop and setup as above, but using a DCGAN.

## 1.5 Build DCGAN

**Task:** Implement the generator of a DCGAN. To allow you to train the DCGAN also with a different dataset later, we change the input resolution from 28x28 as before to 64x64.

### 1.5.1 Generator

```
In [14]: g_c = torch.nn.Sequential(
    #nn.ConvTranspose2d(latent_size, 1024, kernel_size=(4, 4), stride=(1, 1), bias=False),
    nn.ConvTranspose2d(latent_size, 512, kernel_size=(4, 4), stride=(1, 1), bias=False),
    nn.BatchNorm2d(512),
    nn.ReLU(),
    nn.ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=1),
    nn.BatchNorm2d(256),
```

```

        nn.ReLU(),
        nn.ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(),
        nn.ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(),
        nn.ConvTranspose2d(64, 1, kernel_size=(4, 4), stride=(2, 2), padding=1),
        nn.Tanh()
    )

```

### 1.5.2 Discriminator

**Task:** Implement the discriminator of a DCGAN. To allow you to train the DCGAN also with a different dataset later, we change the input resolution from 28x28 as before to 64x64.

```

In [19]: d_c = torch.nn.Sequential(
    nn.Conv2d(1, 64, 4, 2, 1, bias=False),
    nn.LeakyReLU(negative_slope=0.2, inplace=True),
    nn.Conv2d(64, 128, 4, 2, 1, bias=False),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(negative_slope=0.2, inplace=True),
    nn.Conv2d(128, 256, 4, 2, 1, bias=False),
    nn.BatchNorm2d(256),
    nn.LeakyReLU(negative_slope=0.2, inplace=True),
    nn.Conv2d(256, 512, 4, 2, 1, bias=False),
    nn.BatchNorm2d(512),
    nn.LeakyReLU(negative_slope=0.2, inplace=True),
    nn.Conv2d(512, 1, 4, 1, 0, bias=False),
    nn.Sigmoid()
)

```

```

In [21]: test_input = torch.randn((128, 100, 1, 1))
out = g_c(test_input)

```

```

In [22]: out.shape

```

```

Out[22]: torch.Size([128, 1, 64, 64])

```

**Task:** Once again, we have to send the networks to our device and initialize the optimizers. Use the same learning rate as above.

```

In [23]: criterion = nn.BCELoss()

        # send to device
        netG_c = g_c.to(device)
        netD_c = d_c.to(device)

        # optimizer
        d_optimizer_c = Adam(netD_c.parameters(), lr=0.0002, betas=(0.5, 0.999))
        g_optimizer_c = Adam(netG_c.parameters(), lr=0.0002, betas=(0.5, 0.999))

```

### 1.5.3 Transforms

**Task:** To adapt the MNIST images to the new expected input dimension of 64x64, we need to adapt our transforms accordingly. Find the right transform to resize all images to 64x64 and initialize dataset and data loader with the new transforms.

```
In [24]: new_image_size = 64*64
        transform_c = transforms.Compose([
            transforms.Resize((64, 64)),
            #transforms.CenterCrop(image_size),
            transforms.ToTensor(),
            transforms.Normalize([0.5], [0.5])
        ])
        dataset = torchvision.datasets.MNIST(root='.',
                                              train=True,
                                              transform=transform_c,
                                              download=True)
        data_loader = torch.utils.data.DataLoader(dataset=dataset,
                                                batch_size=batch_size,
                                                shuffle=True, drop_last=True)
```

### 1.5.4 Training Loop Reloaded

Now we can build the training loop for the DCGAN. But don't worry, we can mostly copy our code from above. Take care of the changed tensor dimensions, though.

**Task:** Implement the Training loop from above and visualize a batch of fake data after every epoch. How does the training differ from the vanilla GAN training?

```
In [25]: total_step = len(data_loader)
        for epoch in range(num_epochs):
            for i, (images, _) in enumerate(data_loader):
                # Create the labels which are later used as input for the BCE loss
                real_labels = torch.ones(batch_size, 1).to(device)
                fake_labels = torch.zeros(batch_size, 1).to(device)
                images = images.to(device)
                # ===== #
                #           Train the discriminator                      #
                # ===== #
                outputs = d_c(images)
                d_loss_real = criterion(outputs, real_labels)
                real_score = outputs

                # Compute BCELoss using fake images
                z = torch.randn(batch_size, latent_size, 1, 1).to(device)
                fake_images = g_c(z)
                outputs = d_c(fake_images)
                d_loss_fake = criterion(outputs, fake_labels)
                fake_score = outputs
```

```

# Backprop and optimize
d_loss = d_loss_real + d_loss_fake
d_optimizer_c.zero_grad()
g_optimizer_c.zero_grad()
d_loss.backward()
d_optimizer_c.step()

# ===== #
#           Train the generator                      #
# ===== #

# Compute loss with fake images
z = torch.randn(batch_size, latent_size, 1, 1).to(device)
fake_images = g_c(z)
outputs = d_c(fake_images)

# We train G to maximize log(D(G(z)) instead of minimizing log(1-D(G(z)))
g_loss = criterion(outputs, real_labels)

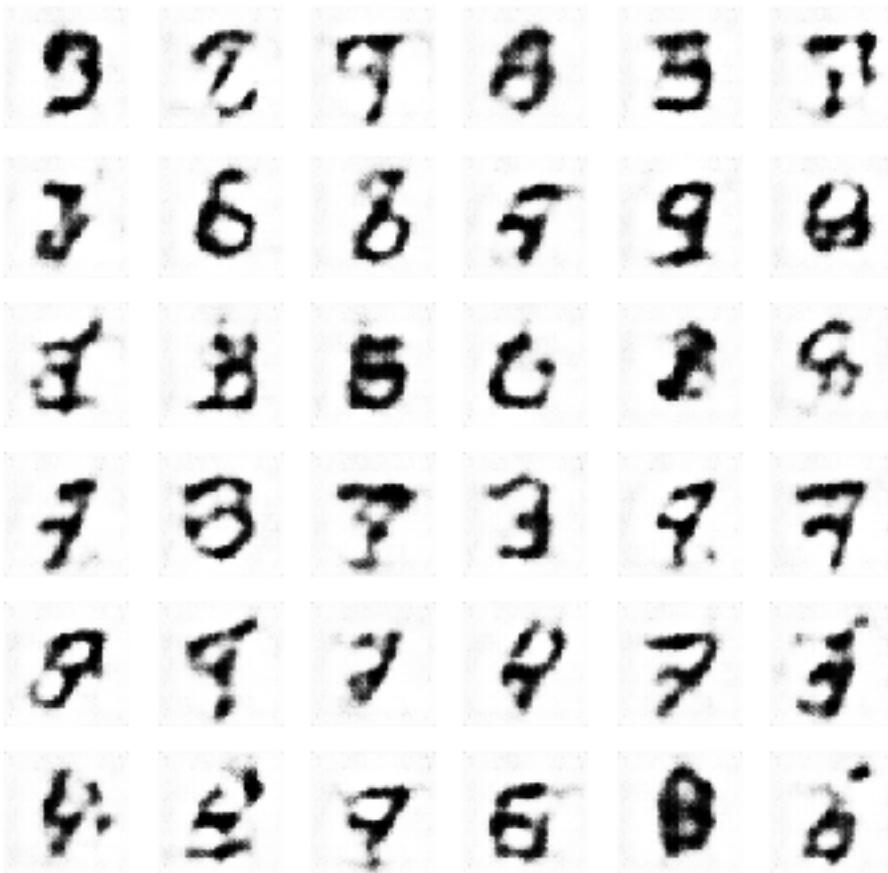
# Backprop and optimize
d_optimizer_c.zero_grad()
g_optimizer_c.zero_grad()
g_loss.backward()
g_optimizer_c.step()

if (i+1) % 200 == 0:
    print('Epoch [{}/{}], Step [{}/{}], d_loss: {:.4f}, g_loss: {:.4f}, D(x): {:.4f}, D(G(z)): {:.4f}'.format(epoch, num_epochs, i+1, total_step, d_loss.item(), g_loss.item(), real_score.mean().item(), fake_score.mean().item()))
    fake_images = fake_images.reshape(fake_images.size(0), 1, 64, 64).cpu().detach().numpy()
    plot_batch(fake_images)

/usr/local/anaconda3/envs/hamlet-pytorch-1-1/lib/python3.7/site-packages/torch/nn/modules/loss.py:187: UserWarning: F.binary_cross_entropy is deprecated in favor of nn.BCELoss. Please use nn.BCELoss instead.
  return F.binary_cross_entropy(input, target, weight=self.weight, reduction=self.reduction)

```

Epoch [0/30], Step [200/468], d\_loss: 0.1821, g\_loss: 4.7172, D(x): 0.86, D(G(z)): 0.02  
Epoch [0/30], Step [400/468], d\_loss: 0.6930, g\_loss: 2.6134, D(x): 0.79, D(G(z)): 0.34



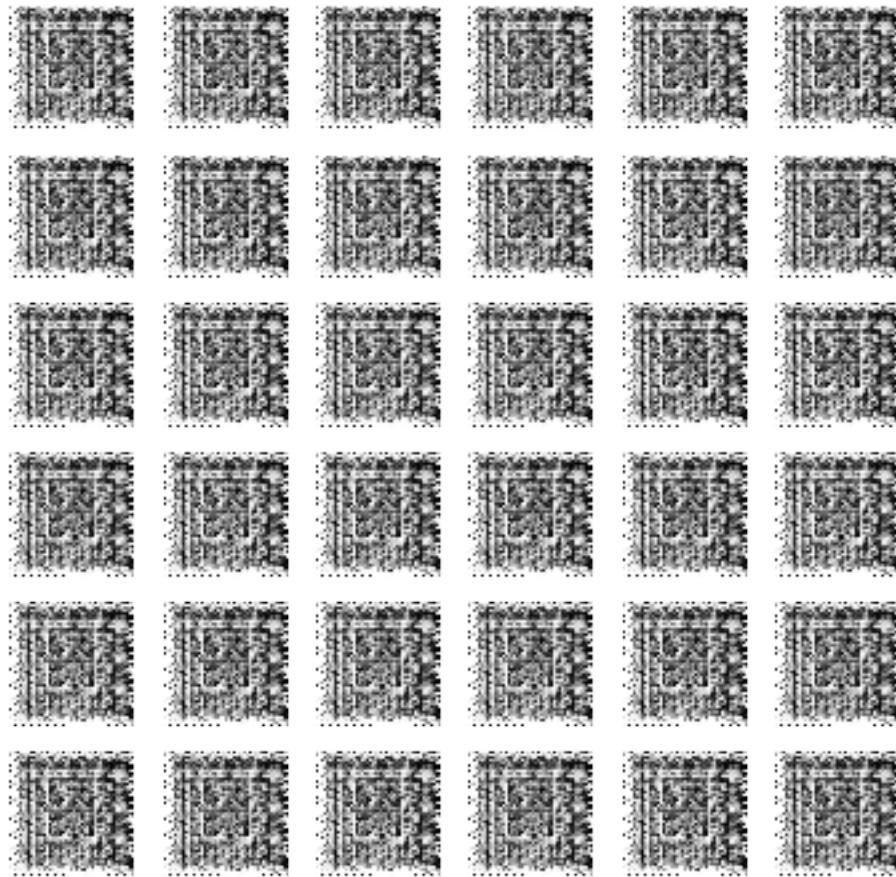
Epoch [1/30], Step [200/468], d\_loss: 0.6200, g\_loss: 2.2121, D(x): 0.80, D(G(z)): 0.32  
Epoch [1/30], Step [400/468], d\_loss: 0.5635, g\_loss: 2.4475, D(x): 0.80, D(G(z)): 0.28

1 9 5 8 0 7  
0 7 8 2 8 0  
3 8 1 4 7 4  
5 7 0 0 1 5  
2 0 9 5 7 3  
3 2 9 3 1 5

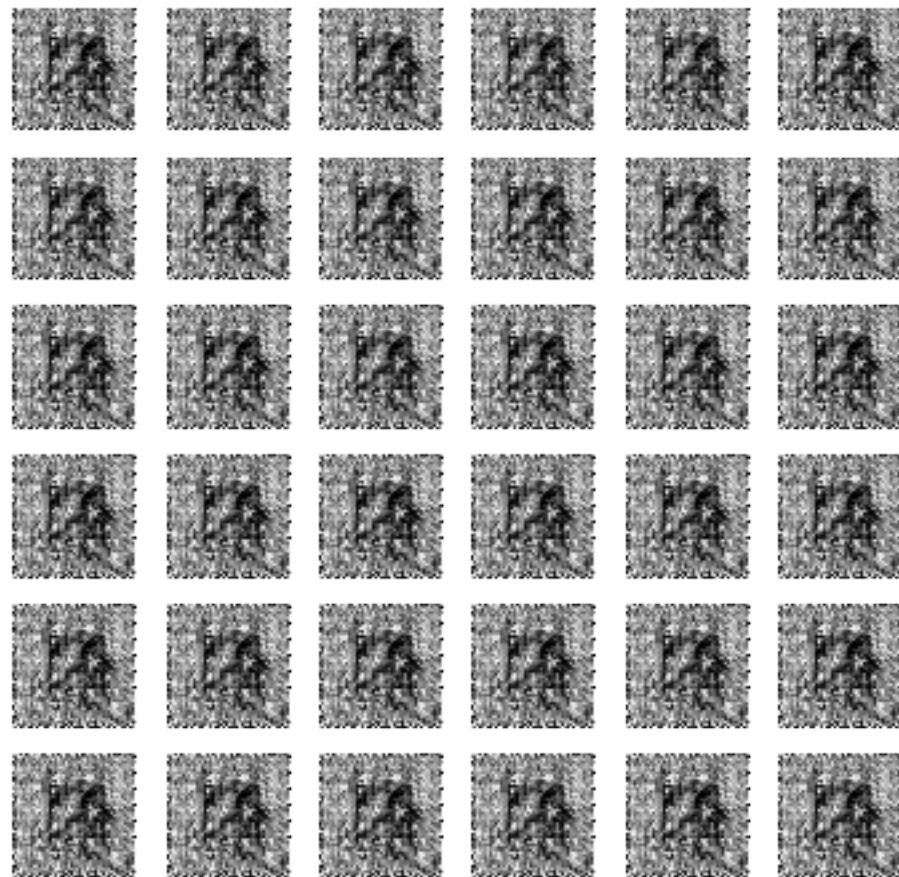
Epoch [2/30], Step [200/468], d\_loss: 0.5096, g\_loss: 3.1736, D(x): 0.87, D(G(z)): 0.29  
Epoch [2/30], Step [400/468], d\_loss: 0.7519, g\_loss: 2.3375, D(x): 0.81, D(G(z)): 0.39



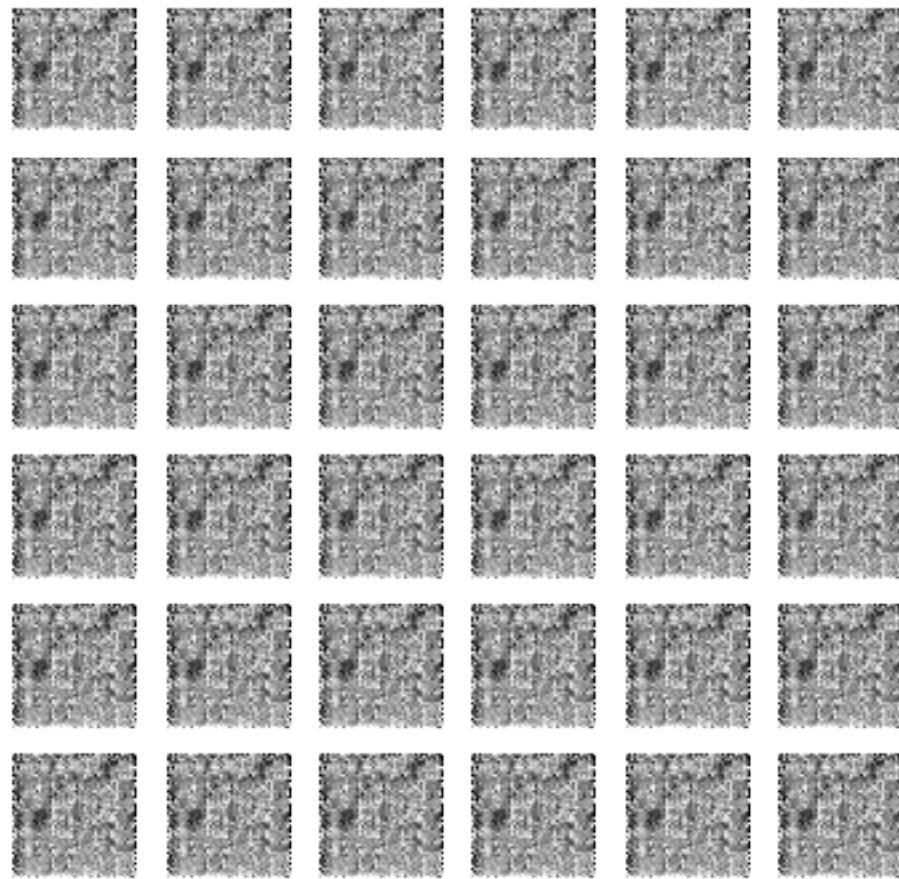
Epoch [3/30], Step [200/468], d\_loss: 0.6667, g\_loss: 1.5133, D(x): 0.76, D(G(z)): 0.29  
Epoch [3/30], Step [400/468], d\_loss: 0.0138, g\_loss: 6.7604, D(x): 1.00, D(G(z)): 0.01



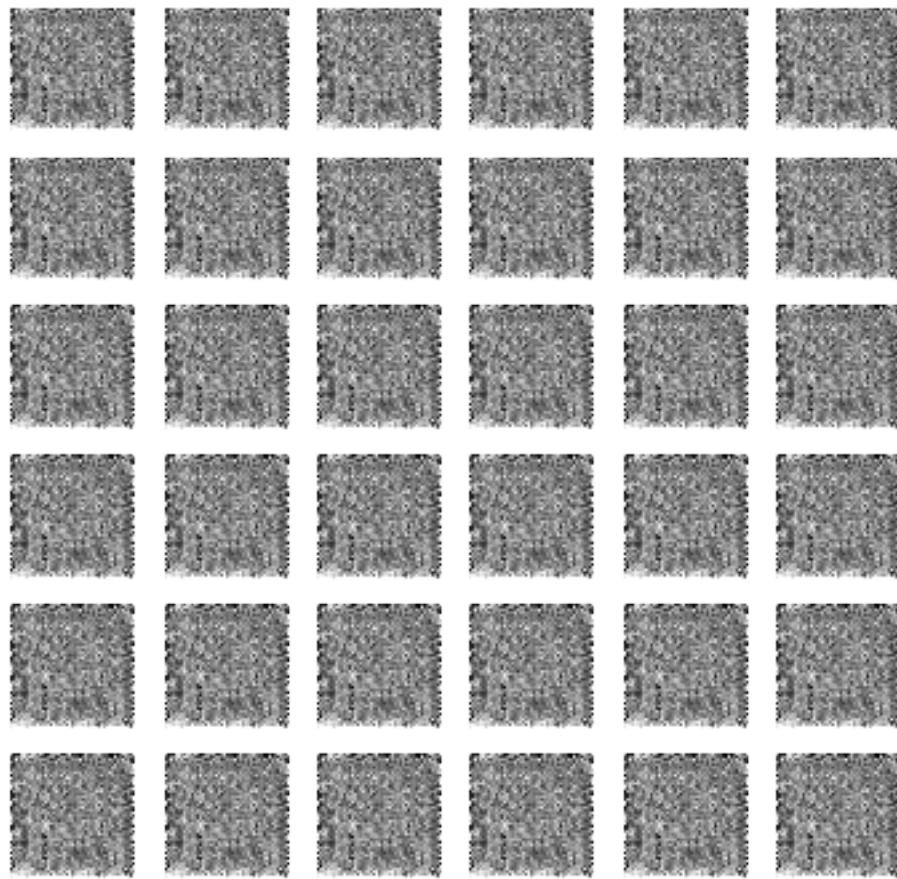
Epoch [4/30], Step [200/468], d\_loss: 0.0015, g\_loss: 7.0911, D(x): 1.00, D(G(z)): 0.00  
Epoch [4/30], Step [400/468], d\_loss: 0.0016, g\_loss: 6.9488, D(x): 1.00, D(G(z)): 0.00



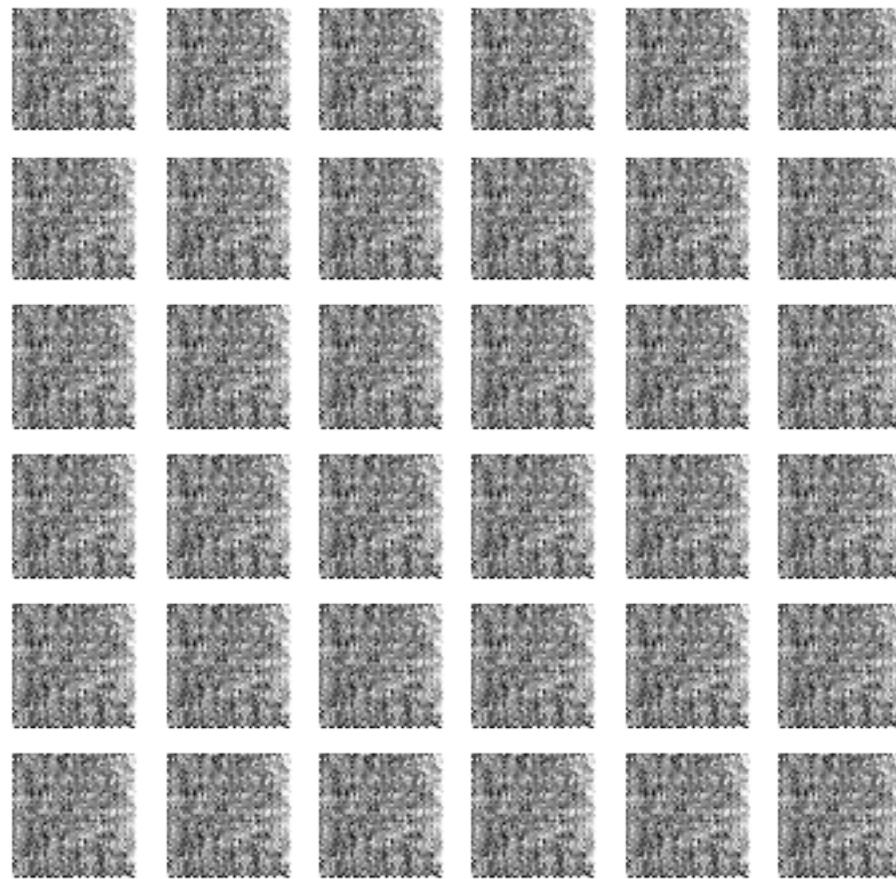
Epoch [5/30], Step [200/468], d\_loss: 0.0003, g\_loss: 8.7054, D(x): 1.00, D(G(z)): 0.00  
Epoch [5/30], Step [400/468], d\_loss: 0.0004, g\_loss: 8.5067, D(x): 1.00, D(G(z)): 0.00



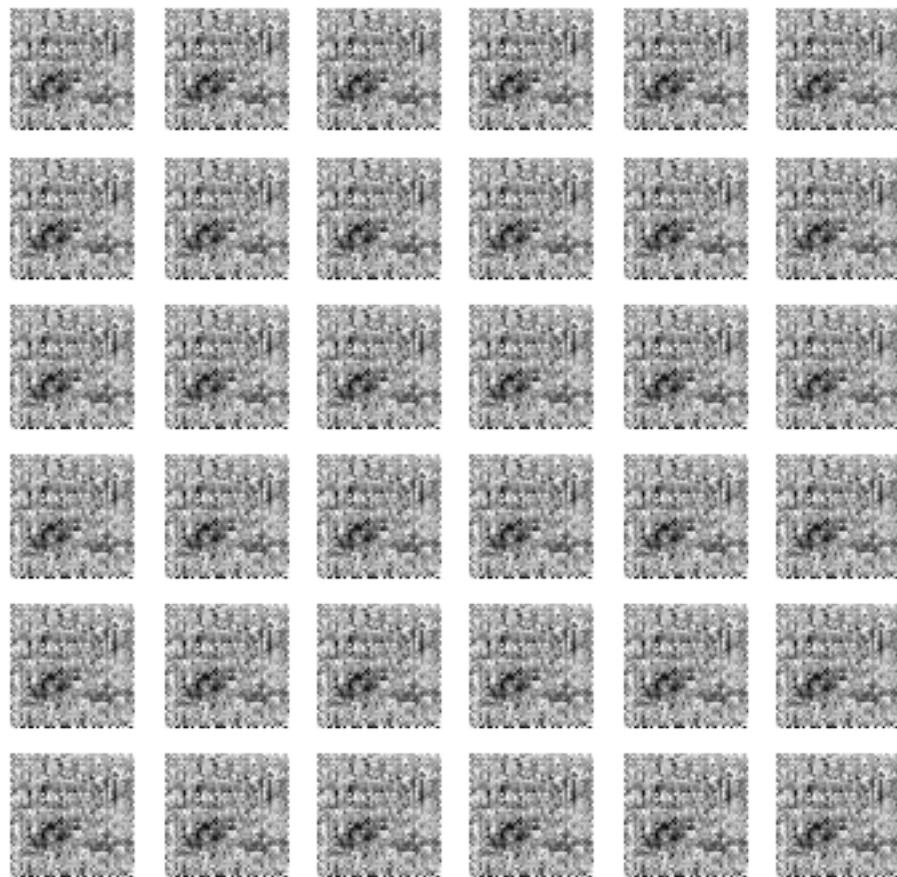
Epoch [6/30], Step [200/468], d\_loss: 0.0002, g\_loss: 8.6291, D(x): 1.00, D(G(z)): 0.00  
Epoch [6/30], Step [400/468], d\_loss: 0.0002, g\_loss: 8.8919, D(x): 1.00, D(G(z)): 0.00



Epoch [7/30], Step [200/468], d\_loss: 0.0002, g\_loss: 8.7816, D(x): 1.00, D(G(z)): 0.00  
Epoch [7/30], Step [400/468], d\_loss: 0.0001, g\_loss: 9.6698, D(x): 1.00, D(G(z)): 0.00



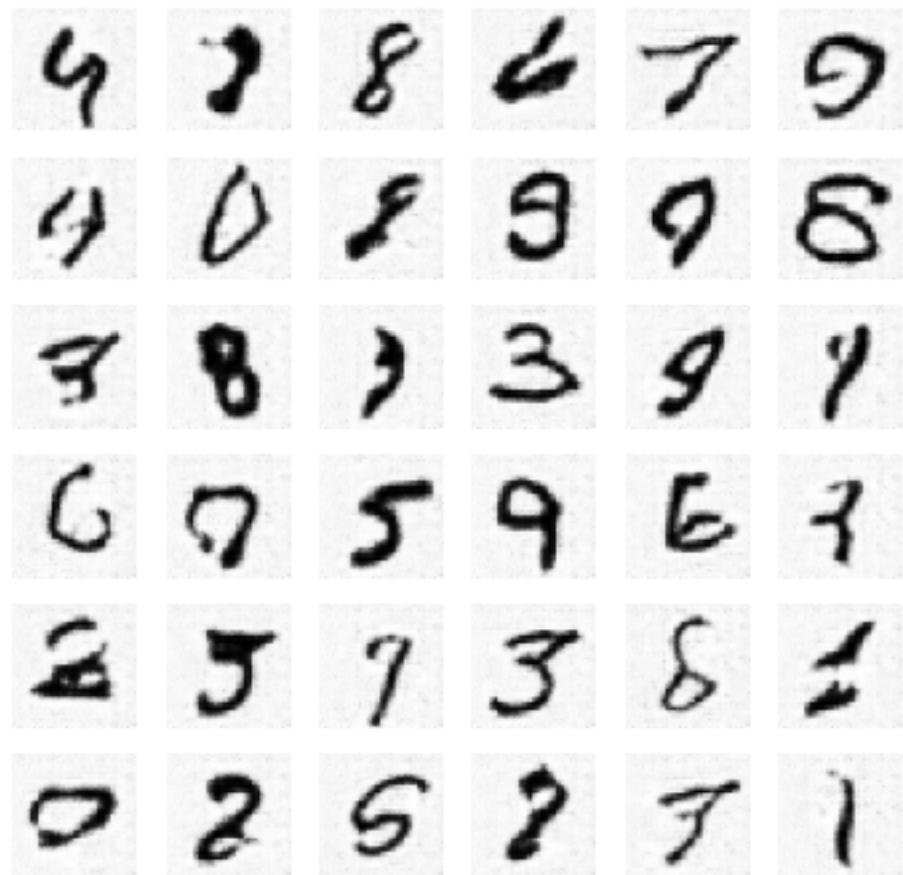
Epoch [8/30], Step [200/468], d\_loss: 0.0001, g\_loss: 10.1828, D(x): 1.00, D(G(z)): 0.00  
Epoch [8/30], Step [400/468], d\_loss: 0.0001, g\_loss: 10.0996, D(x): 1.00, D(G(z)): 0.00



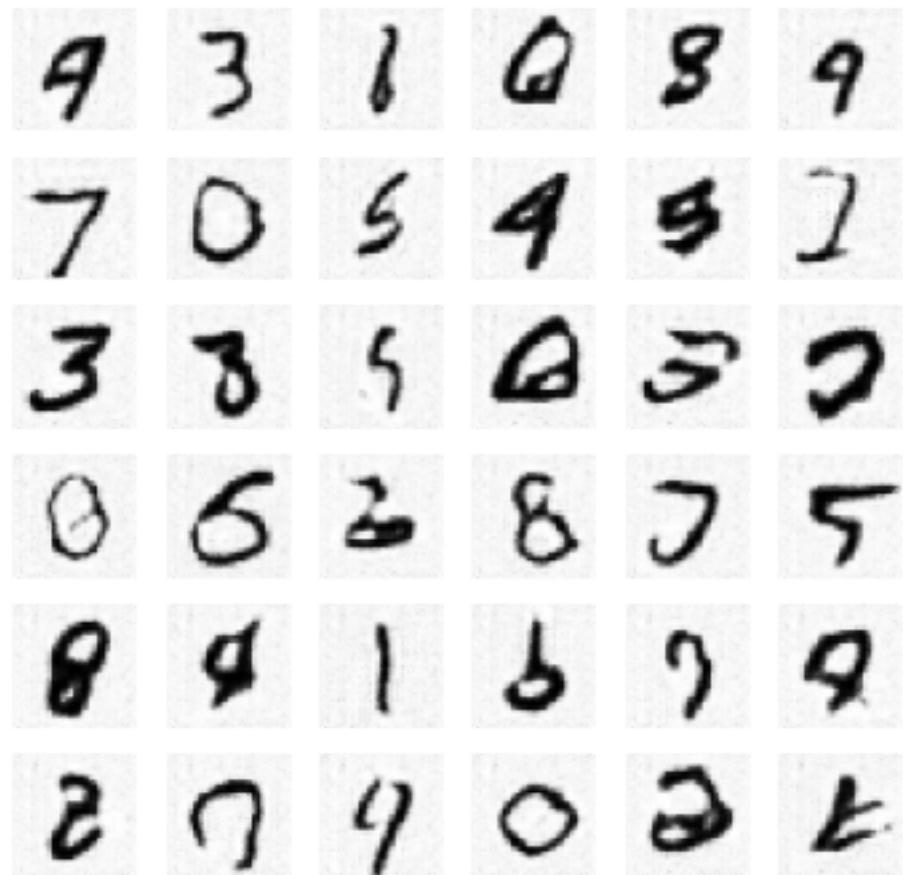
Epoch [9/30], Step [200/468], d\_loss: 0.0710, g\_loss: 7.1816, D(x): 0.94, D(G(z)): 0.00  
Epoch [9/30], Step [400/468], d\_loss: 1.4461, g\_loss: 2.0296, D(x): 0.31, D(G(z)): 0.00



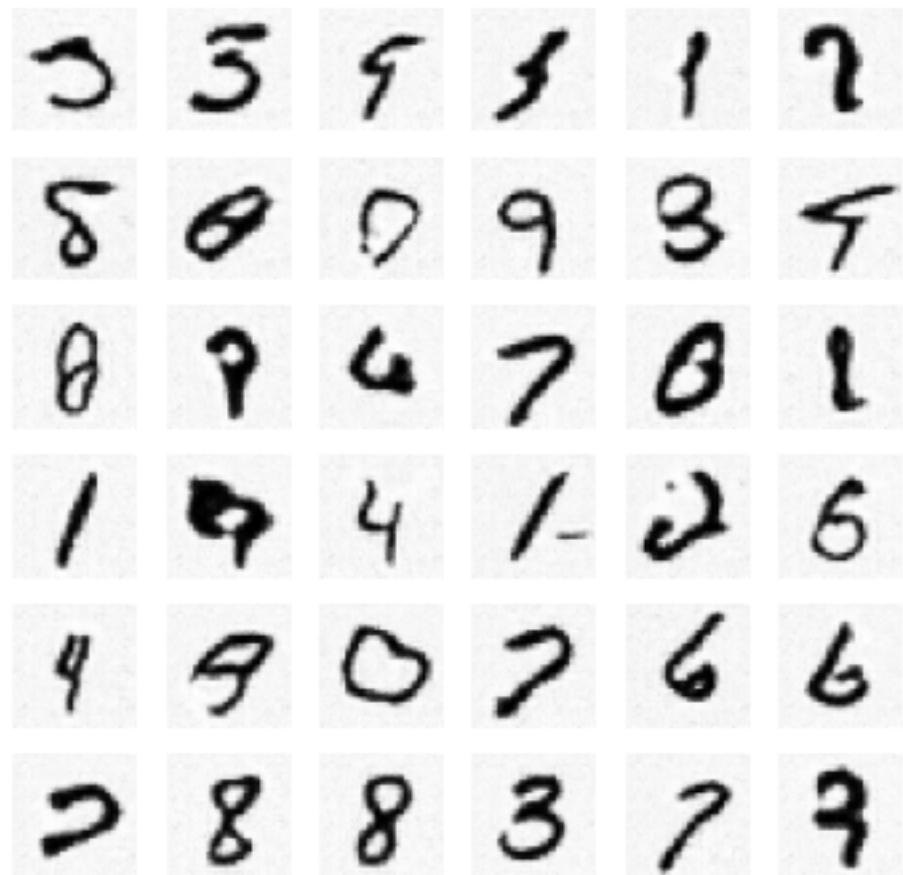
Epoch [10/30], Step [200/468], d\_loss: 0.3965, g\_loss: 2.1246, D(x): 0.77, D(G(z)): 0.11  
Epoch [10/30], Step [400/468], d\_loss: 0.3616, g\_loss: 3.4598, D(x): 0.90, D(G(z)): 0.22



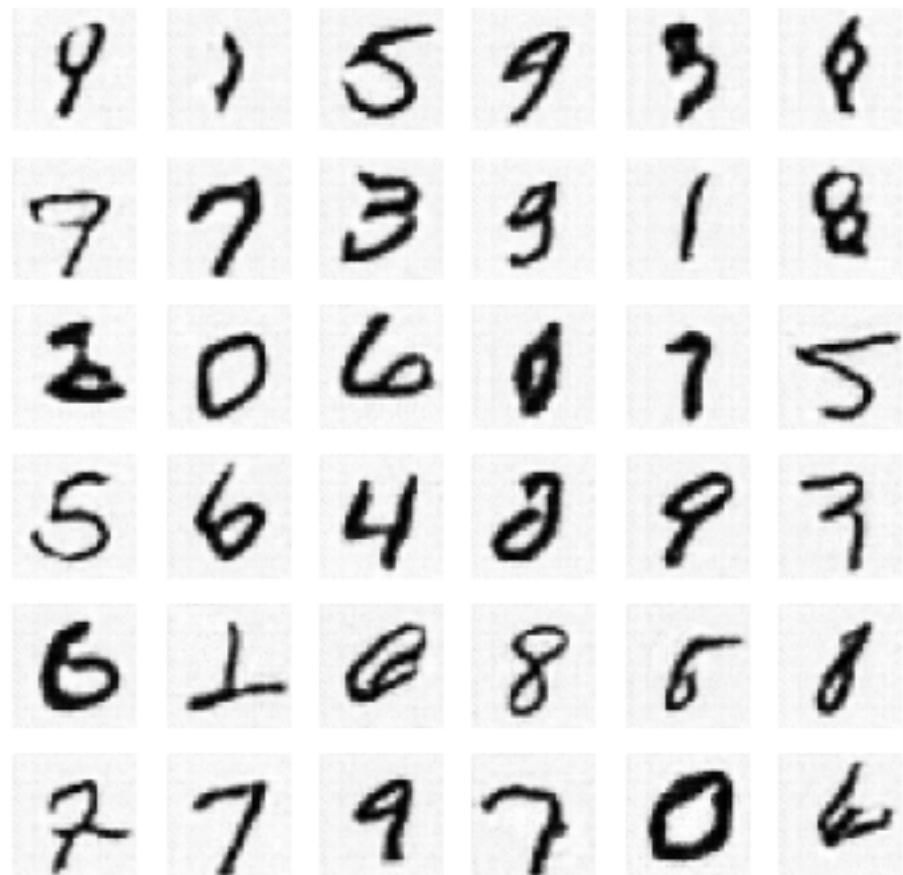
Epoch [11/30], Step [200/468], d\_loss: 1.5872, g\_loss: 0.7308, D(x): 0.27, D(G(z)): 0.02  
Epoch [11/30], Step [400/468], d\_loss: 1.4905, g\_loss: 1.3133, D(x): 0.84, D(G(z)): 0.67



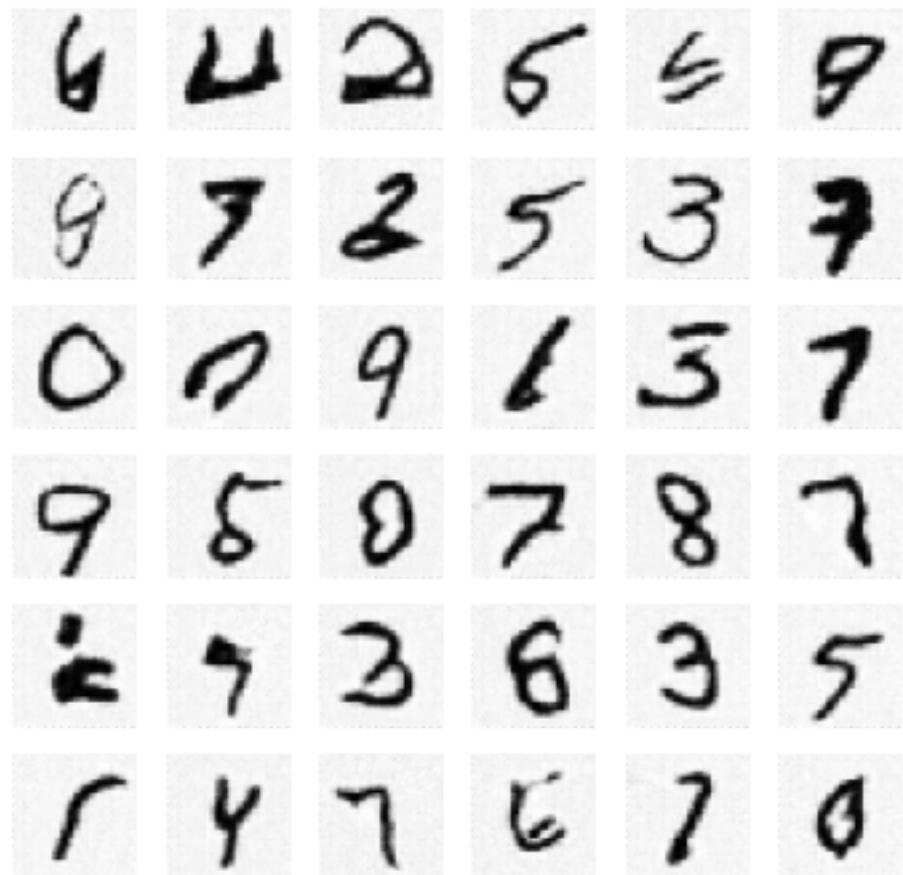
Epoch [12/30], Step [200/468], d\_loss: 0.1648, g\_loss: 2.2819, D(x): 0.89, D(G(z)): 0.04  
Epoch [12/30], Step [400/468], d\_loss: 0.4075, g\_loss: 1.5605, D(x): 0.71, D(G(z)): 0.01



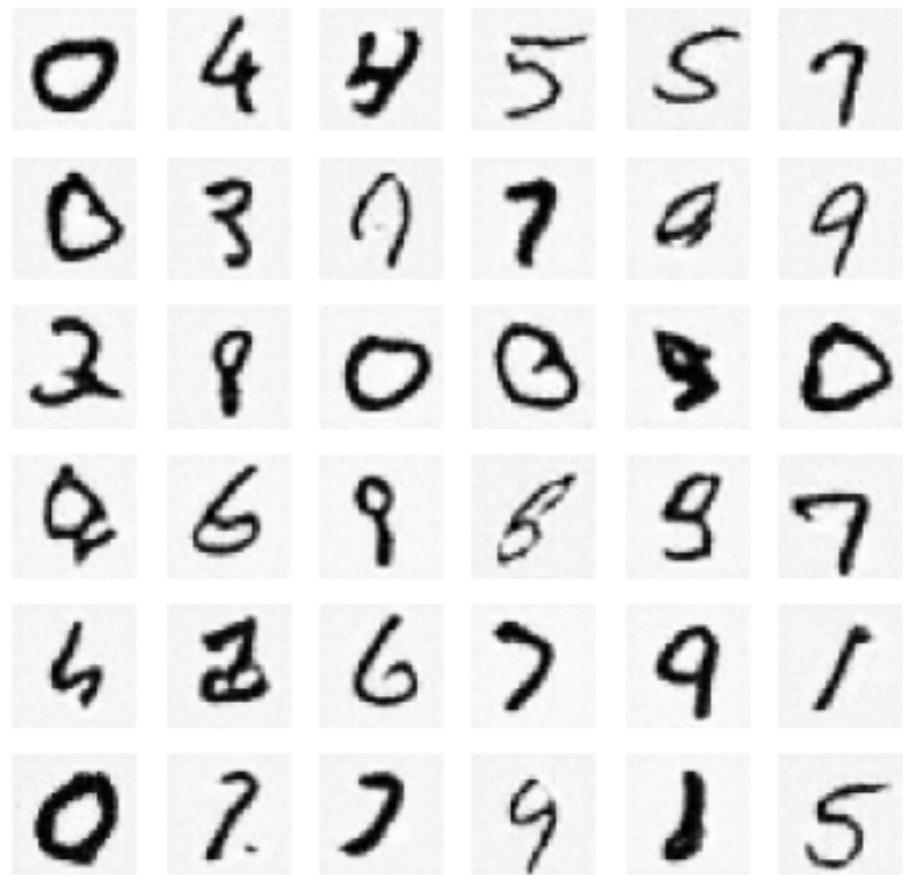
Epoch [13/30], Step [200/468], d\_loss: 0.0768, g\_loss: 3.7021, D(x): 0.97, D(G(z)): 0.04  
Epoch [13/30], Step [400/468], d\_loss: 0.0620, g\_loss: 4.0260, D(x): 0.96, D(G(z)): 0.01



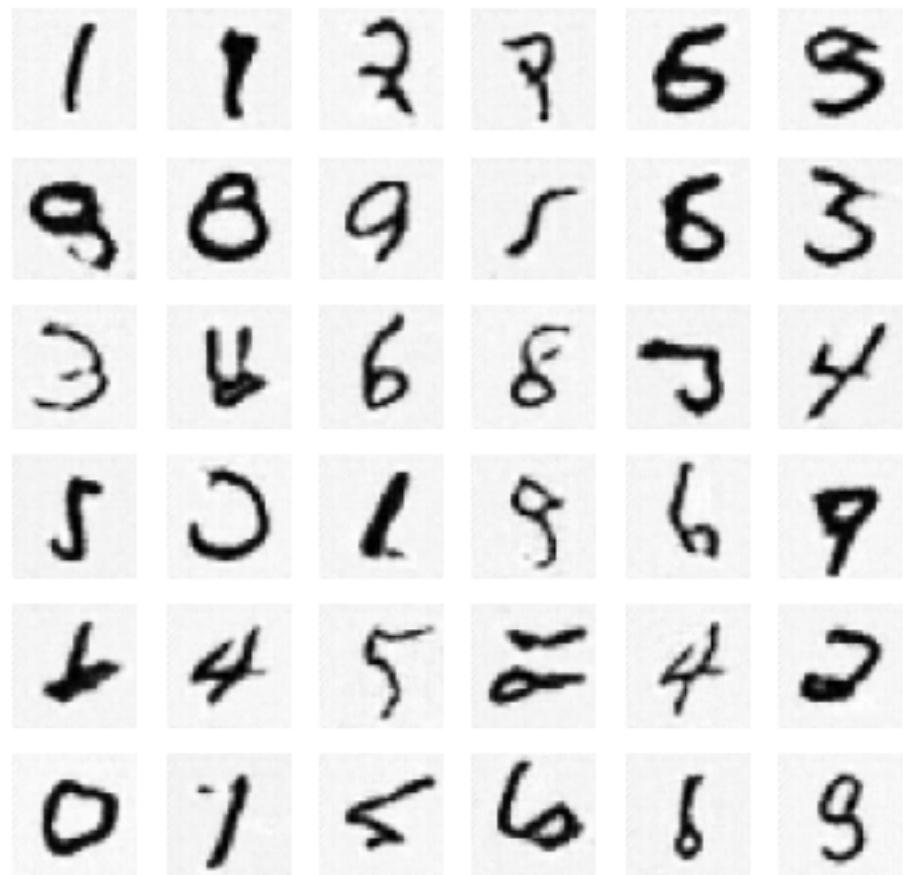
Epoch [14/30], Step [200/468], d\_loss: 1.6406, g\_loss: 6.9093, D(x): 1.00, D(G(z)): 0.74  
Epoch [14/30], Step [400/468], d\_loss: 4.9307, g\_loss: 6.9496, D(x): 1.00, D(G(z)): 0.99



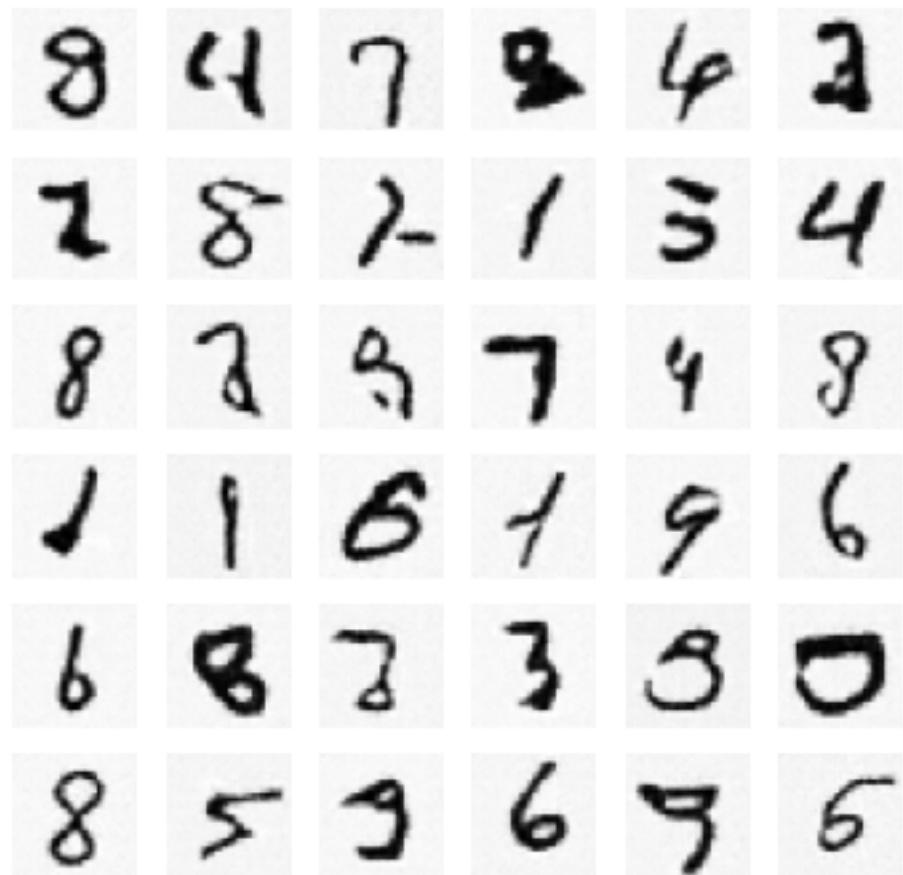
Epoch [15/30], Step [200/468], d\_loss: 0.0896, g\_loss: 3.9697, D(x): 0.95, D(G(z)): 0.03  
Epoch [15/30], Step [400/468], d\_loss: 0.1241, g\_loss: 3.2350, D(x): 0.97, D(G(z)): 0.09



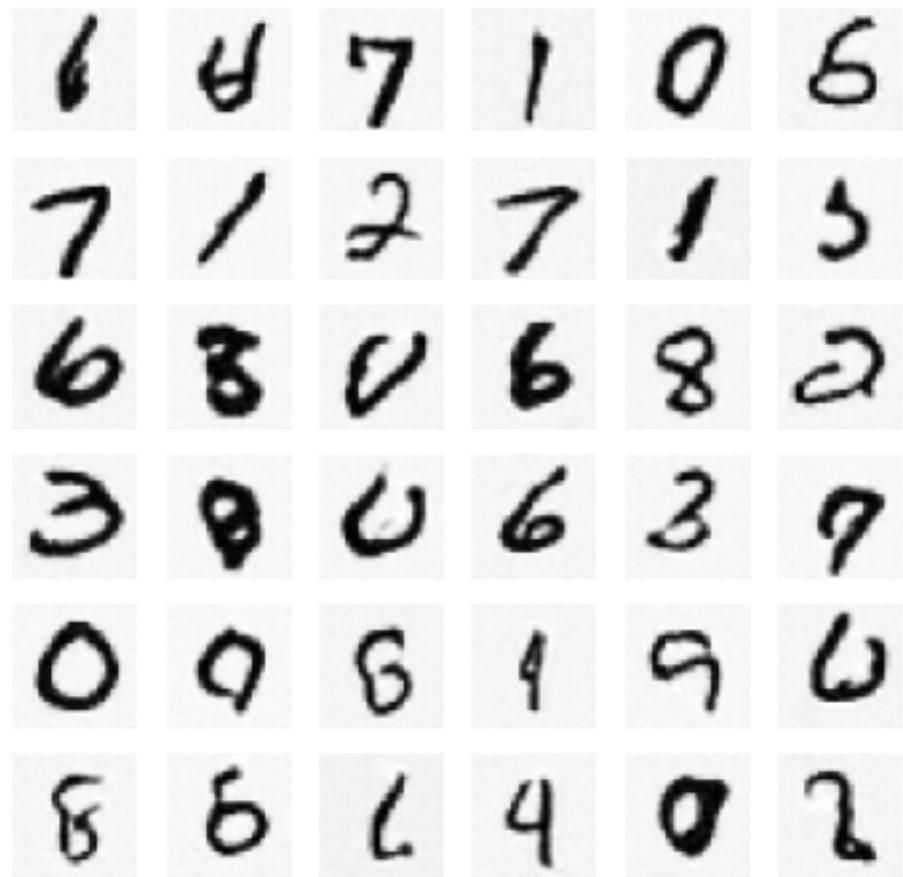
Epoch [16/30], Step [200/468], d\_loss: 0.1439, g\_loss: 3.2796, D(x): 0.89, D(G(z)): 0.02  
Epoch [16/30], Step [400/468], d\_loss: 0.2105, g\_loss: 3.2038, D(x): 0.90, D(G(z)): 0.09



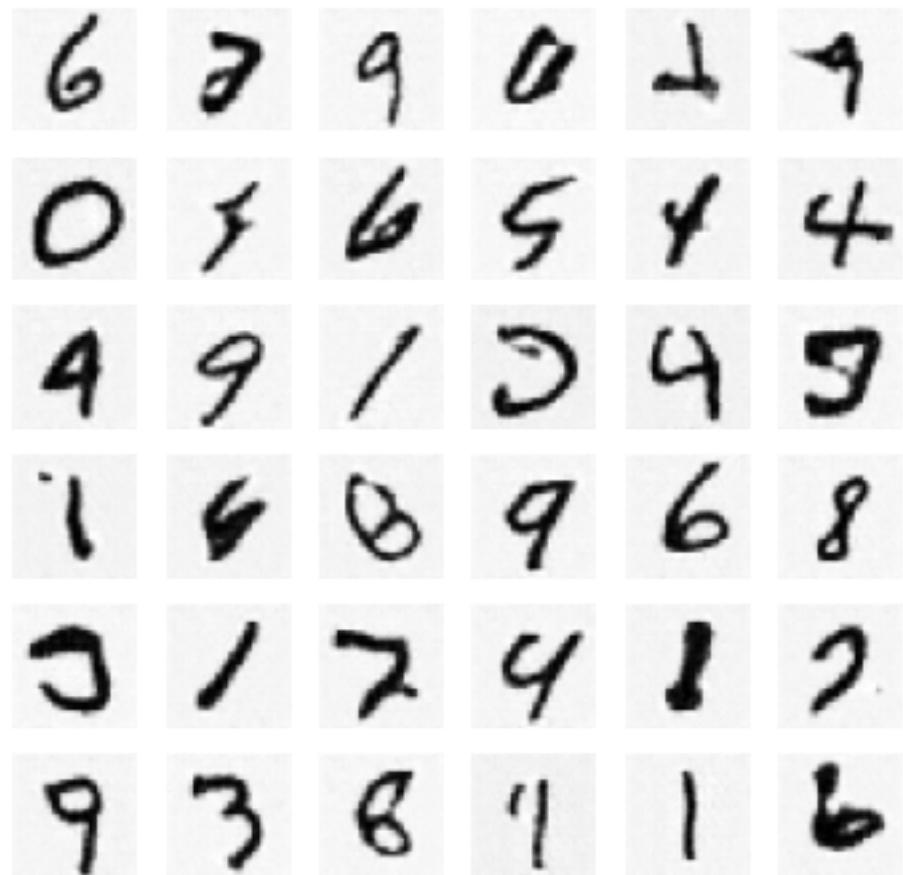
Epoch [17/30], Step [200/468], d\_loss: 0.0258, g\_loss: 4.3833, D(x): 0.99, D(G(z)): 0.02  
Epoch [17/30], Step [400/468], d\_loss: 0.2890, g\_loss: 4.4560, D(x): 0.96, D(G(z)): 0.20



Epoch [18/30], Step [200/468], d\_loss: 1.9407, g\_loss: 4.2014, D(x): 0.99, D(G(z)): 0.78  
Epoch [18/30], Step [400/468], d\_loss: 6.3546, g\_loss: 0.1717, D(x): 0.00, D(G(z)): 0.00



Epoch [19/30], Step [200/468], d\_loss: 0.0346, g\_loss: 4.1575, D(x): 0.98, D(G(z)): 0.02  
Epoch [19/30], Step [400/468], d\_loss: 0.0306, g\_loss: 4.7252, D(x): 0.99, D(G(z)): 0.02



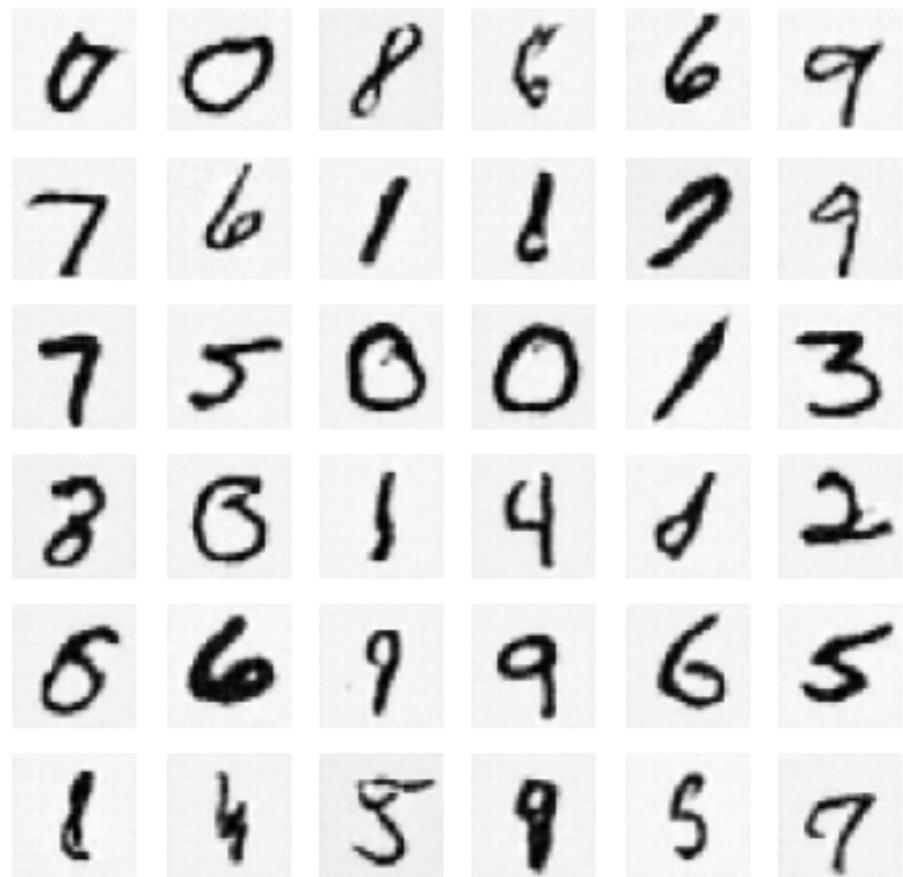
Epoch [20/30], Step [200/468], d\_loss: 0.1457, g\_loss: 3.4130, D(x): 0.96, D(G(z)): 0.09  
Epoch [20/30], Step [400/468], d\_loss: 0.0549, g\_loss: 4.5773, D(x): 0.95, D(G(z)): 0.01



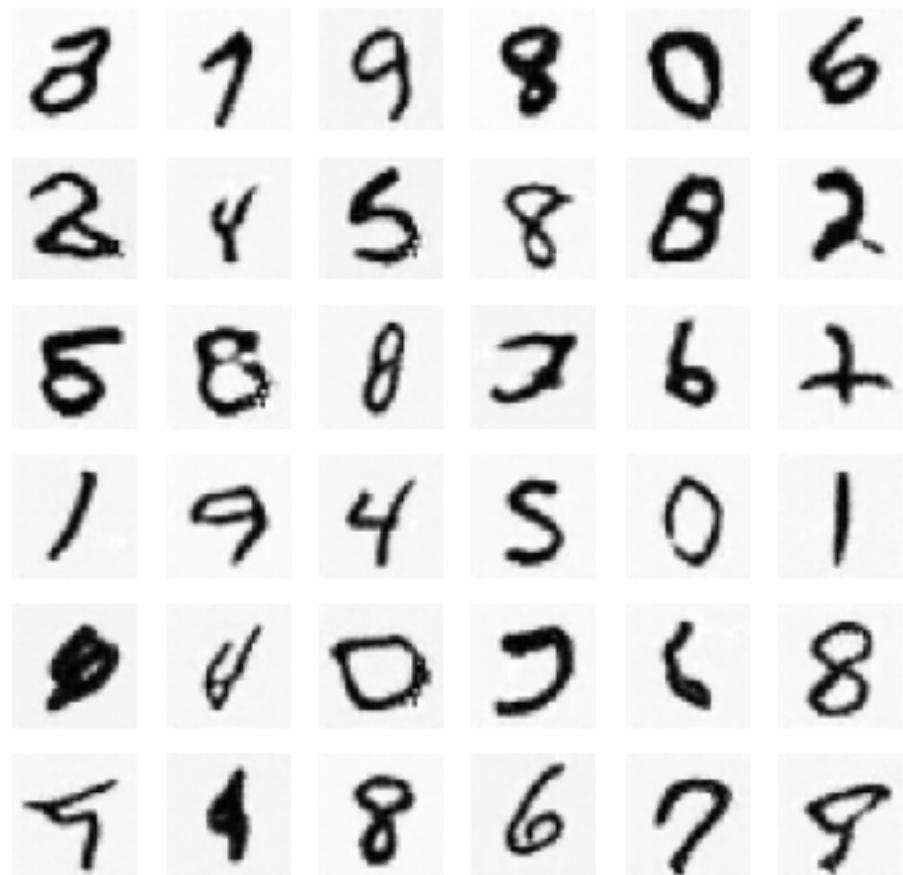
Epoch [21/30], Step [200/468], d\_loss: 0.0332, g\_loss: 4.8634, D(x): 0.98, D(G(z)): 0.01  
Epoch [21/30], Step [400/468], d\_loss: 0.4759, g\_loss: 2.8831, D(x): 0.93, D(G(z)): 0.30



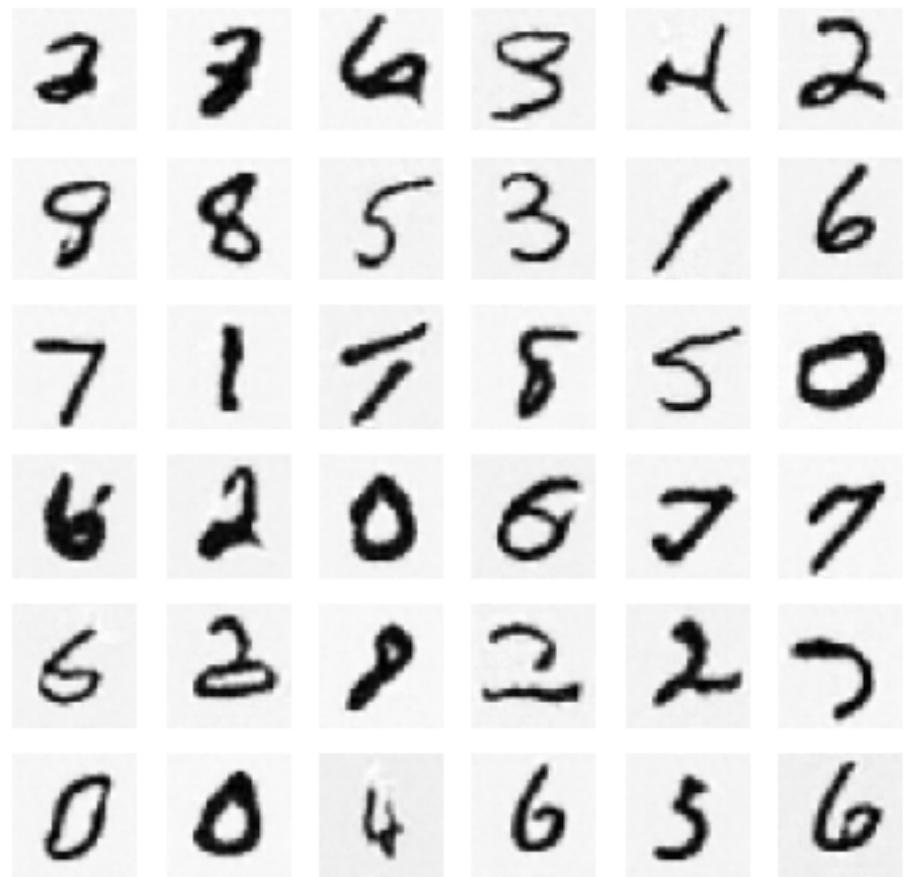
Epoch [22/30], Step [200/468], d\_loss: 0.0345, g\_loss: 4.3618, D(x): 0.98, D(G(z)): 0.01  
Epoch [22/30], Step [400/468], d\_loss: 0.7046, g\_loss: 1.3828, D(x): 0.69, D(G(z)): 0.24



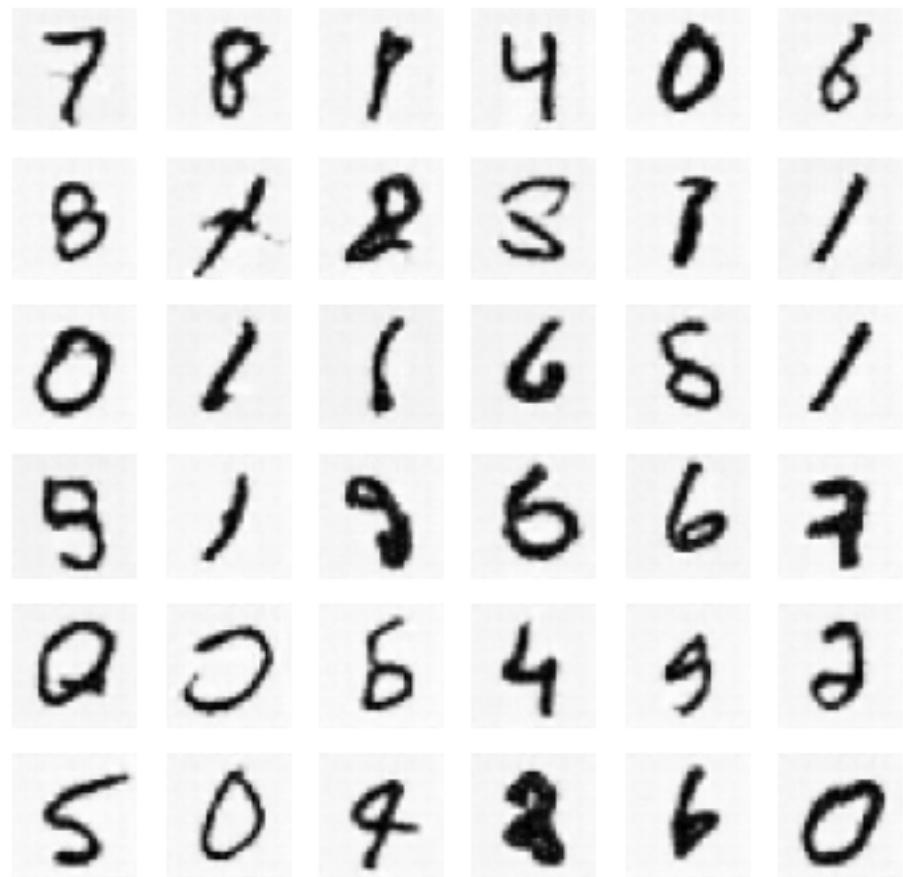
Epoch [23/30], Step [200/468], d\_loss: 0.0296, g\_loss: 5.2937, D(x): 0.98, D(G(z)): 0.01  
Epoch [23/30], Step [400/468], d\_loss: 0.0142, g\_loss: 4.3330, D(x): 0.99, D(G(z)): 0.01



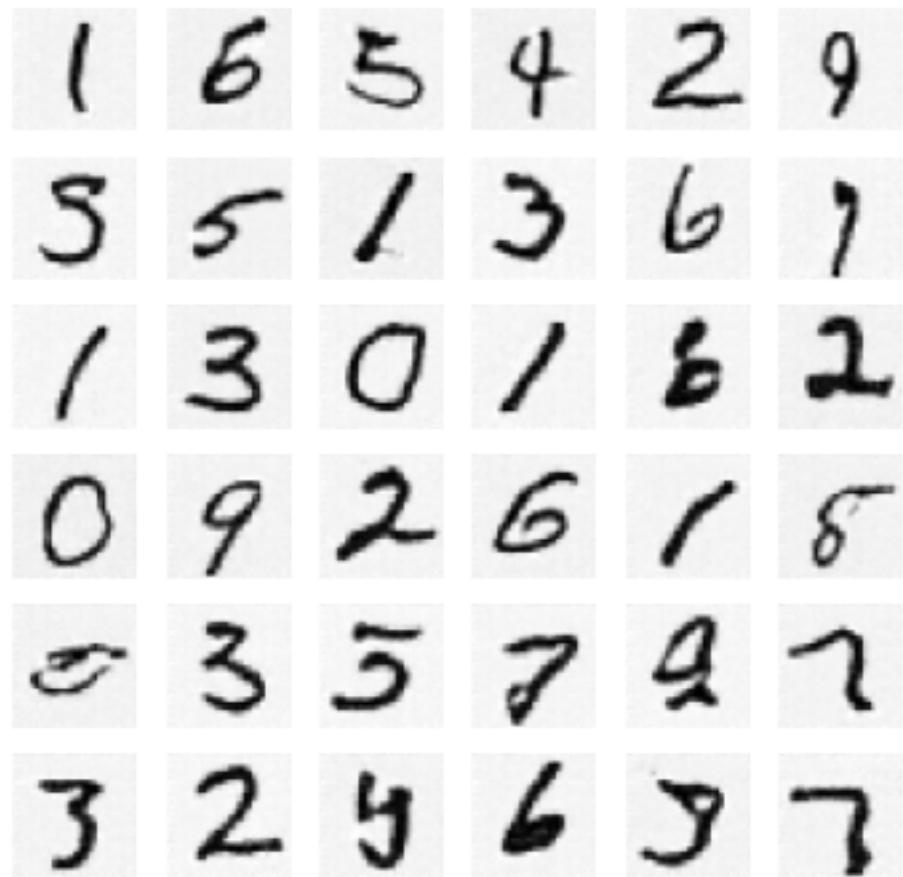
Epoch [24/30], Step [200/468], d\_loss: 1.0245, g\_loss: 1.9638, D(x): 0.82, D(G(z)): 0.49  
Epoch [24/30], Step [400/468], d\_loss: 0.1030, g\_loss: 3.2958, D(x): 0.95, D(G(z)): 0.05



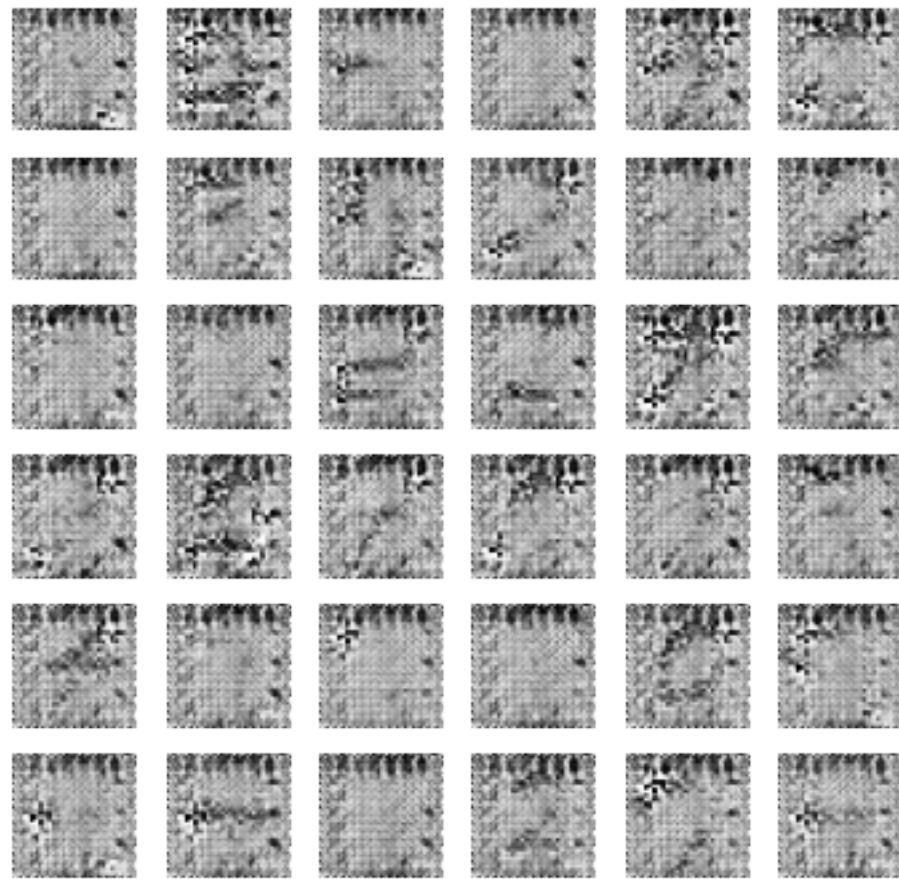
Epoch [25/30], Step [200/468], d\_loss: 0.0761, g\_loss: 3.8846, D(x): 0.97, D(G(z)): 0.05  
Epoch [25/30], Step [400/468], d\_loss: 0.0130, g\_loss: 6.0150, D(x): 0.99, D(G(z)): 0.01



Epoch [26/30], Step [200/468], d\_loss: 1.0556, g\_loss: 1.9765, D(x): 0.85, D(G(z)): 0.52  
Epoch [26/30], Step [400/468], d\_loss: 0.1076, g\_loss: 1.9102, D(x): 0.91, D(G(z)): 0.01



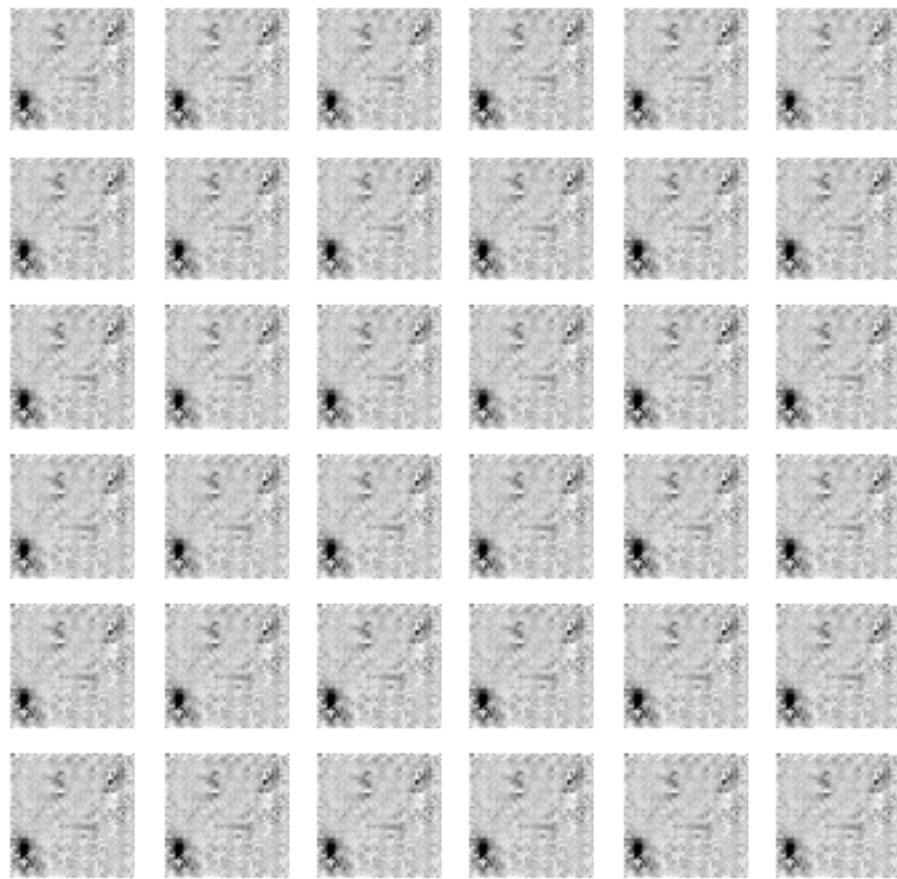
Epoch [27/30], Step [200/468], d\_loss: 0.0257, g\_loss: 4.1618, D(x): 0.99, D(G(z)): 0.01  
Epoch [27/30], Step [400/468], d\_loss: 0.0129, g\_loss: 5.6174, D(x): 0.99, D(G(z)): 0.01



Epoch [28/30], Step [200/468], d\_loss: 0.0039, g\_loss: 7.0930, D(x): 1.00, D(G(z)): 0.00  
Epoch [28/30], Step [400/468], d\_loss: 0.0012, g\_loss: 7.5325, D(x): 1.00, D(G(z)): 0.00



Epoch [29/30], Step [200/468], d\_loss: 0.0006, g\_loss: 7.8921, D(x): 1.00, D(G(z)): 0.00  
Epoch [29/30], Step [400/468], d\_loss: 0.0005, g\_loss: 8.0703, D(x): 1.00, D(G(z)): 0.00



## 1.6 Bonus Task

Congratulations, you've successfully trained a DCGAN that can generate MNIST numbers. If you are brave, it's time to generate some more complex images.

**Task:** Have a look at the datasets that are included in `torchvision` and train a DCGAN with another dataset.

### 1.6.1 Feedback

That's it, we're done

If you have any suggestions on how we could improve this session, please let us know in the following cell. What did you particularly like or dislike? Did you miss any contents?